

C++ now

C++ Type Erasure Demystified

Fedor Pikus

2024

C++ Type Erasure Demystified

Fedor G Pikus
Siemens Fellow

Outline

- What is type erasure
- What does type erasure look like
- Type erasure as a design pattern
- Type erasure as an implementation technique
- How does it work?
- Three ways to implement type erasure in C++
 - Inheritance
 - Static functions
 - V-table
- Performance benchmarks

Outline

- What is type erasure
- What does type erasure look like
- Type erasure as a design pattern
- Type erasure as an implementation technique
- How does it work?
- Three ways to implement type erasure in C++
 - Inheritance
 - Static functions
 - V-table
- Performance benchmarks

What is type erasure?

- Type erasure is a programming technique by which the explicit type information is removed from the program. It is a type of abstraction that ensures that the program does not explicitly depend on some of the data types.

What is type erasure?

- Type erasure is magic.

What is type erasure?

- Type erasure is a programming technique by which the explicit type information is removed from the program. It is a type of abstraction that ensures that the program does not explicitly depend on some of the data types.
- A program is written in a strongly typed language but does not use the actual types. How?
- Why, by abstracting away the type, of course!

Outline

- What is type erasure
- What does type erasure look like
- Type erasure as a design pattern
- Type erasure as an implementation technique
- How does it work?
- Three ways to implement type erasure in C++
 - Inheritance
 - Static functions
 - V-table
- Performance benchmarks

How does type erasure look like?

- How does a program with explicit types look like?

```
{  
    std::unique_ptr<int> p(new int(0));  
}
```

- Creates and deletes an integer
- Deletion is not explicitly visible
 - Done by `std::default_delete` which calls operator `delete`
 - How does a program with **explicit** types look like?

How does type erasure look like?

- How does a program with explicit types look like?

```
class MyHeap {  
    ...  
    void* allocate(size_t size);  
    void deallocate(void* p);  
};  
  
void* operator new(size_t size, MyHeap* heap) {  
    return heap->allocate(size);  
}
```

- Custom operator new(size_t, MyHeap*) for MyHeap
 - no operator delete with arguments

How does type erasure look like?

- How does a program with explicit types look like?

```
struct MyDeleter {  
    ...  
    template <typename T> void operator()(T* p);  
};  
{  
    MyHeap heap;  
    std::unique_ptr<int, MyDeleter> p(new(&heap) int(0),  
                                     MyDeleter(&heap));  
}
```

- Creates and deletes an integer, allocation from heap

How does type erasure look like?

- How does a program with explicit types look like?

```
class MyDeleter {  
    MyHeap* heap_;  
public:  
    MyDeleter(MyHeap* heap) : heap_(heap) {}  
    template <typename T> void operator()(T* p) {  
        p->~T();  
        heap_>deallocate(p);  
    }  
};
```

- No-throw movable (or copyable)

Show me the explicit types

- Types are explicitly present in the program

```
MyHeap heap;
```

```
std::unique_ptr<int, MyDeleter> p(new(&heap) int(0),  
                                MyDeleter(&heap));
```

```
std::unique_ptr<int> q(new int(0));
```

```
p = std::move(q); // Error: p and q are different types
```

```
std::unique_ptr<int> p(new(&heap) int(0), // default_delete is  
                    MyDeleter(&heap)); // not MyDeleter
```

- Unique pointers to different types are different types – of course
- Unique pointers to the same type but with different deleters are different types too (we can deduce deleter type from the pointer type)

Where is type erasure already?

- How is shared pointer different from unique pointer?

```
std::unique_ptr<int> p(new int(0));  
std::shared_ptr<int> q(new int(0));
```

- Now with custom deleter:

```
MyHeap heap;  
std::unique_ptr<int, MyDeleter> p(new(&heap) int(0),  
                                   MyDeleter(&heap));  
std::shared_ptr<int> q(new(&heap) int(0),  
                       MyDeleter(&heap));
```

- Where is the deleter type? Erased!
- We cannot deduce deleter type from the pointer type

Is the erased type gone?

- Shared pointers with different deleters have the same type

```
{  
    std::shared_ptr<int> p(new int(0));  
    MyHeap heap;  
    std::shared_ptr<int> q(new(&heap) int(0), MyDeleter(&heap));  
    q = p; // OK, same type  
} // Proper deleters are called!
```

- But each shared pointer invokes the correct deleter
- Erased types are not explicitly visible in the program (no `decltype` in your code depends on the erased type)
- Actions that depend on these types are performed correctly

Outline

- What is type erasure
- What does type erasure look like
- Type erasure as a design pattern
- Type erasure as an implementation technique
- How does it work?
- Three ways to implement type erasure in C++
 - Inheritance
 - Static functions
 - V-table
- Performance benchmarks

General type erasure

```
std::function<size_t(const std::string&)> f;
```

- Any function that can be called with a string and returns a `size_t`

```
size_t f1(const std::string& s) { return s.capacity(); }
```

```
f = f1;
```

```
f = [=](const std::string& s) { return s.find(c); };
```

```
f = &std::string::size;
```

- `f1`, lambda, and member function all have different types
 - `f` has only one type but can store any of these callable objects
- Type erasure is an abstraction for multiple implementations that provide the same behavior (the relevant behavior is what matters)

Type erasure as a design pattern

- What problem does type erasure solve
 - The code expects certain behavior
 - The code is written in terms of an abstraction that provides this behavior
 - Many concrete types can implement this behavior
 - All properties of these types that are not relevant to the behavior are erased
 - Starting with the name of the type
- Type erasure separates the interface from the implementation
 - So does inheritance, but type erasure does not require common base class
 - Type erasure is non-intrusive
 - External polymorphism (types do not have to be designed for it)

Outline

- What is type erasure
- What does type erasure look like
- Type erasure as a design pattern
- Type erasure as an implementation technique
- How does it work?
- Three ways to implement type erasure in C++
 - Inheritance
 - Static functions
 - V-table
- Performance benchmarks

Type erasure for decoupling dependencies

- Type erasure can be an implementation technique

```
class Network {  
    void send(const char* data);  
    void receive(const char* buffer);  
};
```

- Network is used by our entire application
 - but one small part needs encryption
 - and another small part has bad network and must use error correction

Type erasure for decoupling dependencies

- Packets may need additional processing

```
class Network {  
    bool needs_processing_  
    void send(const char* data) {  
        if (needs_processing_) apply_processing(data);  
    }  
};
```

- Network is used by our entire application
- All of which now depends on the processing code
- Processors may be of different types

Type erasure for decoupling dependencies

- Type erasure offers a simple solution

```
class Network {  
    std::function<const char*(const char*)> processor =  
        [](const char* c){ return c; };  
    void send(const char* data) {  
        data = processor(data);  
    }  
    template <typename F> void SetProcessor(F&& f) {  
        processor = std::forward<F>(f);  
    }  
};
```

Does not depend on processor type!

One and only mention of processor type

Processor type erased

Type erasure for decoupling dependencies

- Type erasure here implements Strategy pattern
 - Implementation of a particular behavior can be chosen at run time

```
Network N;  
N.SetProcessor(  
    [](const char* s){  
        char* c;  
        ... process the input ...;  
        return c;  
    }  
);
```

Outline

- What is type erasure
- What does type erasure look like
- Type erasure as a design pattern
- Type erasure as an implementation technique
- How does it work?
- Three ways to implement type erasure in C++
 - Inheritance
 - Static functions
 - V-table
- Performance benchmarks

How does the type erasure work?

- You already know how:

```
void qsort(void *base, size_t nmem, size_t size,  
          int (*compare)(const void *, const void *));
```

No mention of specific types

```
int less(const void* a, const void* b) {  
    return *(const int*)a - *(const int*)b;  
}
```

Type information recovered!

```
int a[10] = { 1, 10, 2, 9, 3, 8, 4, 7, 5, 0 };  
qsort(a, 10, sizeof(int), less);
```

Type of less() erased here

- The code depends only on the relevant type properties:
 - size and how to compare types

How does the type erasure work?

- Type erasure in C:

```
void qsort(void *base, size_t nmem, size_t size,  
          int (*compare)(const void *, const void *));
```

- The general code does not depend on the type to sort
- The code depends only on the relevant type properties:
 - size and how to compare types
- All interfaces are generic – no type information

How does the type erasure work?

- Type erasure in C:

```
int a[10] = { 1, 10, 2, 9, 3, 8, 4, 7, 5, 0 };  
qsort(a, 10, sizeof(int), less);
```

- At the call site, the specific types are known
- They may be used to compute some properties (often size)
- All other type information is erased
 - From this point forward, we execute the code that has no knowledge of the specific type it was called with

How does the type erasure work?

- Type erasure in C:

```
int less(const void* a, const void* b) {  
    return *(const int*)a - *(const int*)b;  
}
```

- Type must be recovered at some point
 - where the type-specific actions take place
- Type reification (recovery) is manual in C
 - No compile-time or run-time error detection
- C++ helps with [only] that!
 - C++ automates type reification and makes it correct by construction

The mechanism of type erasure

- The general code does not depend on the erased type
 - Type properties like size are sometimes used
- The call site is the last place where the actual type is known
- Type is reified when the type-specific action must be performed
- Type is hidden in the code of the function that performs this action
- The function is invoked through a type-agnostic interface
 - The type-dependent code converts from abstract to concrete type
- In C, the type-dependent code is written manually
- In C++, we can make the compiler generate the correct code
 - That's “all” C++ adds to type erasure

Type erase implementations


- Three main implementations
 - Using inheritance
 - Using static functions
 - Using v_table
- All done using the shared pointer as the example
- Focus on the deleter, not the shared ownership
- Same as type-erased unique pointer for our purposes
- Owning type-erased smart pointer: `smartptr`

Outline

- What is type erasure
- What does type erasure look like
- Type erasure as a design pattern
- Type erasure as an implementation technique
- How does it work?
- Three ways to implement type erasure in C++
 - Inheritance
 - Static functions
 - V-table
- Performance benchmarks

Type erasure implementation #1

```
template <typename T> class smartptr_te {  
    public:  
        template <typename Deleter> smartptr_te(T* p, Deleter d) :  
            p_(p), ??? {}  
        ~smartptr_te() { ??? delete p using d ??? }  
        T* operator->() { return p_; }  
        const T* operator->() const { return p_; }  
    private:  
        T* p_;  
        ??? something about deleter ???  
};
```



Type erasure implementation #1

```
template <typename T> class smartptr_te {
```

- Only one template parameter – no deleter in the type of smartptr
 - No way to deduce deleter type from smartptr_te type

Type erasure implementation #1

```
template <typename T> class smartptr_te {
```

- Only one template parameter – no deleter in the type of smartptr
- public:

```
template <typename Deleter> smartptr_te(T* p, Deleter d) :  
    p_(p), ??? {}
```

- This constructor is the last place where the type Deleter is known
 - Deleter is not a part of the smartptr_te type
- From this point on, the type is erased and the code is generic
- The constructor must generate some Deleter-specific code
 - And hook it up to the generic call in the destructor

Type erasure implementation #1 - Inheritance

- Erased type Deleter is hidden in a polymorphic derived class:

```
struct destroy_base {  
    virtual void operator()(void*) = 0;  
    virtual ~deleter_base() {}  
};  
template <typename Deleter>  
struct destroy : public destroy_base {  
    destroy (Deleter d) : d_(d) {}  
    void operator()(void* p) override { d_(static_cast<T*>(p)); }  
    Deleter d_;  
};
```

abstract base
(void* only)

derived class template

must use void* to match base

Type erasure implementation #1 - Inheritance

- Erased type Deleter is hidden in a polymorphic derived class

```
template <typename T> class smartptr_te {  
    struct destroy_base { ... };  
    template <typename Deleter>  
    struct destroy : public destroy_base { ... };  
public:  
    template <typename Deleter> smartptr_te(T* p, Deleter d) :  
        p_(p), d_(new destroy<Deleter>(d)) {}  
    ~smartptr_te() { (*d_)(p_); delete d_; }  
    destroy_base* d_;  
};
```

Deleter type known here

type "stored"

no Deleter - type erased

Type erasure implementation #1 - Inheritance

- The constructor of the smart pointer is where we know the type:

```
template <typename Deleter> smartptr_te(T* p, Deleter d)
```

- In the constructor, we create the code that knows the right type:

```
destroy_base d_ = new destroy<Deleter>(d)
```

- The reifying code has type-agnostic interface:

```
template <typename Deleter> struct destroy ... {  
    void operator()(void* p) override { d_(static_cast<T*>(p)); }  
};
```

- This type-agnostic interface is called from the generic code:

```
~smartptr_te() { (*d_)(p_); }
```

Type erasure implementation #1 - Inheritance

- What about default deleter?
 - Bad: leave deleter pointer null, check, and call `std::default_delete`
 - Good: default-initialize the deleter

```
smartptr_te(T* p) : p_(p) {}  
destroy_base d_ = new destroy<std::default_delete<T>>(std::default_delete<T>{});
```

- The rest of the code keeps the same logic!

Type erasure implementation #1 - Inheritance

- Type-erased class does not depend on the erased type
 - `std::shared_ptr<T>`, `std::function<F>`, `std::any`
- The constructor is a template and deduces the type to be erased
 - For the smart pointer, it's the deleter
- The constructor creates a derived object with the override such that:
 - The body of the function uses erased type and is correct by construction
 - The interface of the function is type-agnostic
- The derived object is accessed through the base pointer
- If a default value for the erased type is allowed, the base pointer is default-initialized with the default action

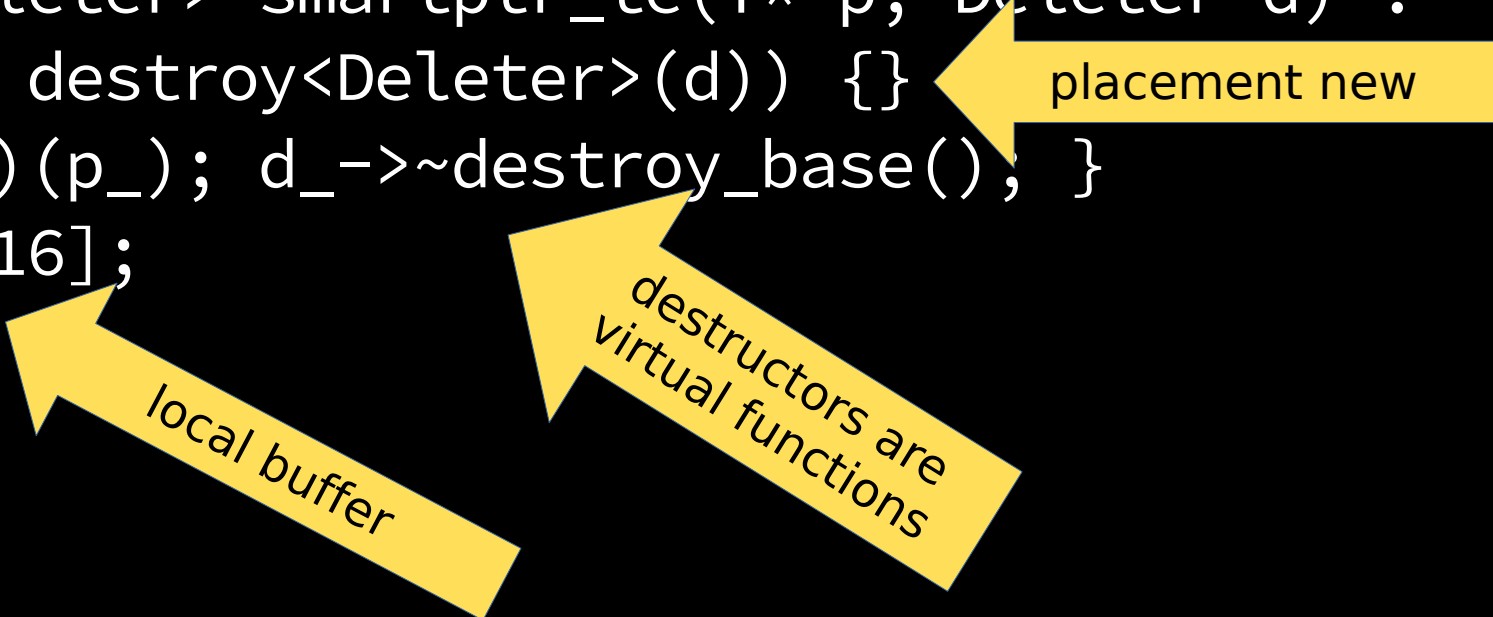
Type erasure implementation #1 - Inheritance

- How to do other common operations on type-erased objects?
- Copying: the destroy hierarchy needs a virtual `clone()` function
- Moving: transfer the deleter to the new object
- Comparison:
 - For smart pointers, often only addresses are compared
 - In general, need another virtual function to compare deleters
- In general, if we need to support an operation on type-erased objects that is affected by the erased type, we have to add a virtual function to the base class and specific overrides to the derived class template
 - Each of these virtual functions needs to reify the type

Type erasure is slow!

- Our implementation of type erasure has a glaring inefficiency: memory is allocated when the type is erased (`new destroy<Deleter>`)
- The common solution is the local buffer optimization:

```
template <typename T> class smartptr_te {  
    template <typename Deleter> smartptr_te(T* p, Deleter d) :  
        p_(p), d_(new(buf_) destroy<Deleter>(d)) {}  
    ~smartptr_te() { (*d_)(p_); d_>~destroy_base(); }  
    alignas(8) char buf_[16];  
    destroy_base* d_;  
};
```



Type erasure is fast but [may be] broken

- Local buffer size does not depend on the erased type `Deleter`
 - The whole point of type erasure is that `smartptr_te` does not depend on it
- When the constructor is called, `Deleter` may or may not fit into `buf_`
- What happens if the erased type does not fit into the local buffer?
 - 1) The implementation switches to dynamic memory allocations
 - 2) `static_assert` in the compiler
- `std::function` uses local buffer and option 1
- Many high-performance implementations use option 2
 - Make the buffer size a template parameter of the class

Type erasure with local buffers

- Local buffer optimization is often used with type erasure
- Avoids dynamic allocations if the type fits into the buffer
- Incurs slight overhead when dynamic allocation is still done
- Design decision: allow all types or only small enough types?
 - Enforced local buffer makes “slow path” a compile-time error
- “Moving” objects with local buffers often becomes copying
 - Size is small, so copy is cheap
 - Copy operations may throw when move is noexcept
- Design decision: restrict the optimization to noexcept-copyable types?

Type erasure using inheritance

- At the point where the type is erased, the compiler instantiates a class template that depends on the erased type
 - Often the constructor of the type-erased class
- All template instantiations inherit from the common base
 - The polymorphic interface is type-agnostic (void*)
- Template generates correct-by-construction member functions that reify the erased type, usually through casts
 - Erased type is hidden in the generated code
- Both primary (deleter) and secondary (copy, move, compare) operations are implemented through virtual overrides

Outline

- What is type erasure
- What does type erasure look like
- Type erasure as a design pattern
- Type erasure as an implementation technique
- How does it work?
- Three ways to implement type erasure in C++
 - Inheritance
 - Static functions
 - V-table
- Performance benchmarks

Type erasure implementation #2

- This technique is much closer to the C type erasure:

```
void reify_func(void* p) {  
    TE* q = static_cast<TE*>(p); ... do work on type TE ...  
}
```

- C++ helps to generate correct-by-construction function:

```
template <typename TE> void reify_func(void* p) {  
    TE* q = static_cast<TE*>(p); ... do work on type TE ...  
}
```

- The signature is type-agnostic – any instantiation can be assigned to

```
void(*) (void*) fp = reify_func<MyDeleter>;
```

Type erasure implementation #2

```
template <typename T> class smartptr_te_static {  
    void(*)(T*, void*) destroy_;  
    template <typename Deleter>  
    static void invoke_destroy(T* p, void* d) {  
        (*static_cast<Deleter*>(d))(p);  
    }  
public:  
    template <typename Deleter>  
    smartptr_te_static(T* p, Deleter d)  
        : p_(p), destroy_(invoke_destroy<Deleter>) { ... }  
};
```

function pointer

no Deleter here

reify template

Deleter used here

type is erased

Type erasure implementation #2

```
template <typename T> class smartptr_te_static {  
    alignas(8) char buf_[8];  
public:  
    template <typename Deleter>  
    smartptr_te_static(T* p, Deleter d)  
        : p_(p), destroy_(invoke_destroy<Deleter>) {  
        ::new (static_cast<void*>(buf_)) Deleter(d);  
    }  
    ~smartptr_te_static() { this->destroy_(p_, buf_); }  
};
```


local buffer

deleter saved

deleter called

Type erasure implementation #2

```
template <typename T> class smartptr_te_static {  
    public:  
    template <typename Deleter>  
    smartptr_te_static(T* p, Deleter d)  
        : p_(p), destroy_(invoke_destroy<Deleter>) {  
        ::new (static_cast<void*>(buf_)) Deleter(d);  
        static_assert(sizeof(Deleter) <= sizeof(buf_));  
        ... also trivially destructible, copyable, etc.  
    }  
};
```



Type erasure implementation #2

- Static functions generated by the template are “code only”
- Only the deleter needs to be stored (not a composite object)
- Local buffer optimization avoids memory allocation costs (same as before)
- Dynamically allocated buffers can be used for larger types
- The downside: how do we copy or destroy the deleter?
 - 1) Limit to trivially destructible/copyable types (often OK in practice)
 - 2) Add another static function to destroy the type-erased deleter
 - And another one for copying...
- This implementation is fast but gets bloated if many operations are abstracted via type erasure

Type erasure using static functions

- At the point where the type is erased, the compiler instantiates a static function template that depends on the erased type
 - Erased type is hidden in the generated code
- All template instantiations have the same signature
 - The signature of the function is type-agnostic (void*)
- Template instantiation is assigned to a function pointer
- Objects with state are stored in local or dynamic memory
- Type-erased code is executed by an indirect function call
- For each supported operation, a function pointer is needed

Outline

- What is type erasure
- What does type erasure look like
- Type erasure as a design pattern
- Type erasure as an implementation technique
- How does it work?
- Three ways to implement type erasure in C++
 - Inheritance
 - Static functions
 - V-table
- Performance benchmarks

Type erasure implementation #3

- This implementation is similar to option 2 (static function) but supports multiple operations without a function pointer for each one

```
template <typename T> class smartptr_te_vtable {  
    struct vtable_t {  
        using destroy_t = void(*) (T*, void*);  
        using destructor_t = void(*) (void*);  
        destroy_t destroy_;  
        destructor_t destructor_;  
    };  
    const vtable_t* vtable_ = nullptr;  
    ...  
};
```

type-agnostic signatures

function pointers

incoming: memory allocation?

Type erasure implementation #3

```
template <typename T> class smartptr_te_vtable {  
    const vtable_t* vtable_ = nullptr;  
    template <typename Deleter>  
    constexpr static vtable_t vtable = { ... };  
public:  
    template <typename Deleter>  
    smartptr_te_vtable(T* p, Deleter d)  
        : p_(p), vtable_(&vtable<Deleter>) {  
    }  
};
```

static template variable

not a template type

no memory allocation

Type erasure implementation #3

```
template <typename T> class smartptr_te_vtable {  
    const vtable_t* vtable_ = nullptr;  
    template <typename Deleter>  
    constexpr static vtable_t vtable = { ... };  
    ...  
    vtable_(&vtable<Deleter>)  
};
```


static template variable

the tricky bit

- Instantiating `vtable` on each `Deleter` type creates a static variable
- Template static variables do not need definitions in `.C` files
- Class `smartptr_te_vtable<T>` has many static variables all named `vtable`

Type erasure implementation #3


```
template <typename T> class smartptr_te_vtable {  
    template <typename Deleter>  
    constexpr static vtable_t vtable = {  
        smartptr_te_vtable::template destroy<Deleter>,  
        smartptr_te_vtable::template destructor<Deleter>  
    };  
    template <typename Deleter>  
    static void destroy(T* p, void* d) {  
        (*static_cast<Deleter*>(d))(p);  
    }  
};
```



```
    struct vtable_t {  
        using destroy_t = void(*) (T*, void*);  
        using destructor_t = void(*) (void*);  
        destroy_t destroy_;  
        destructor_t destructor_;  
    };
```


Type erasure implementation #3

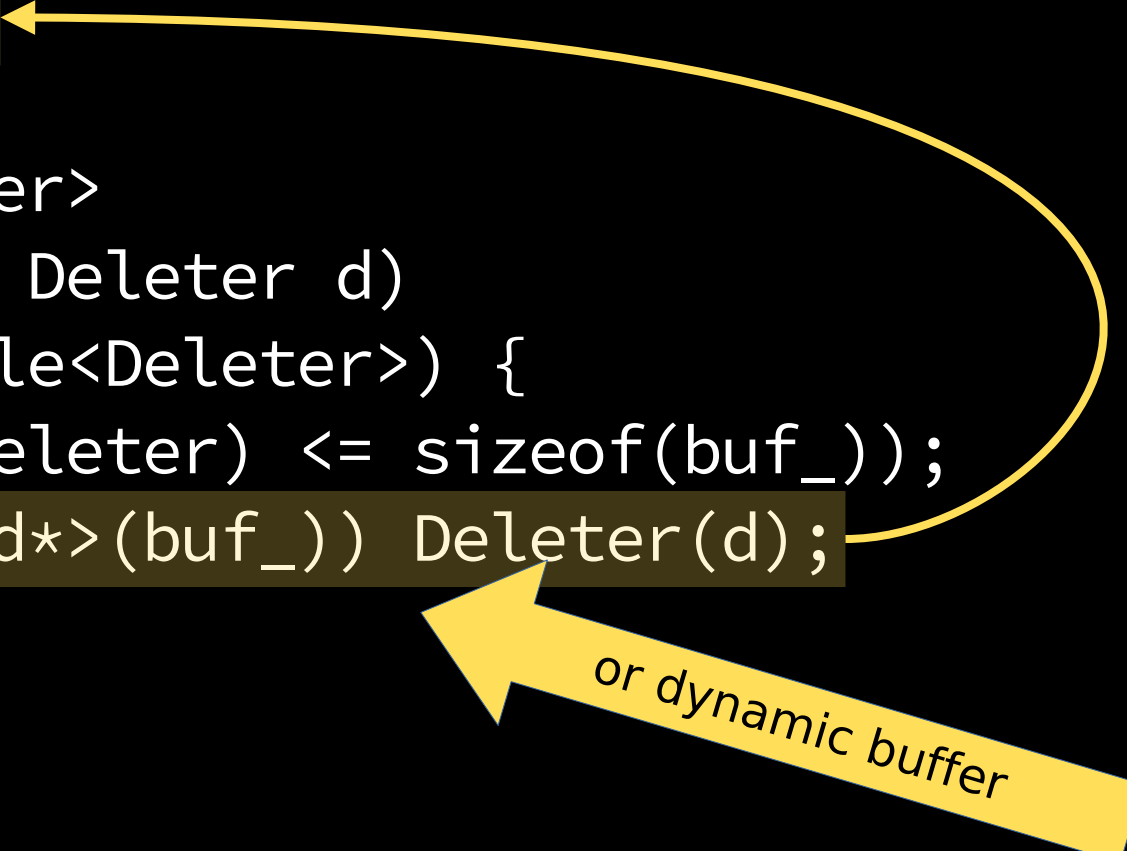
```
template <typename T> class smartptr_te_vtable {  
    template <typename Deleter>  
    constexpr static vtable_t vtable = {  
        smartptr_te_vtable::template destroy<Deleter>,  
        smartptr_te_vtable::template destructor<Deleter>  
    };  
    template <typename Deleter>  
    static void destructor(void* d) {  
        static_cast<Deleter*>(d)->~Deleter();  
    }  
};
```



```
struct vtable_t {  
    using destroy_t = void(*)(T*, void*);  
    using destructor_t = void(*)(void*);  
    destroy_t destroy_;  
    destructor_t destructor_;  
};
```

Type erasure implementation #3

```
template <typename T> class smartptr_te_vtable {  
    const vtable_t* vtable_ = nullptr;  
    alignas(8) char buf_[8];  
public:  
    template <typename Deleter>  
    smartptr_te_vtable(T* p, Deleter d)  
        : p_(p), vtable_(&vtable<Deleter>) {  
        static_assert(sizeof(Deleter) <= sizeof(buf_));  
        ::new (static_cast<void*>(buf_)) Deleter(d);  
    }  
};
```



Type erasure implementation #3

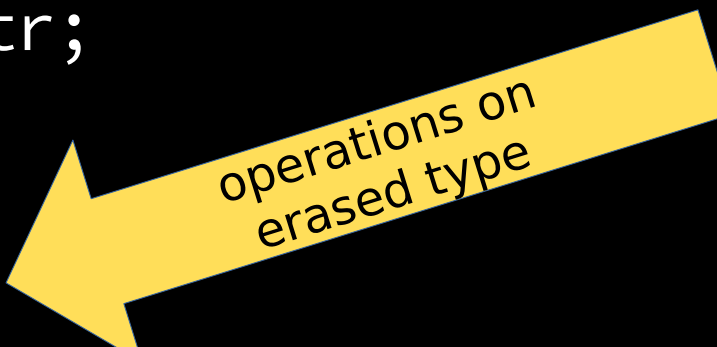
```
template <typename Deleter>
smartptr_te_vtable(T* p, Deleter d)
    : p_(p), vtable_(&vtable<Deleter>) {
    static_assert(sizeof(Deleter) <= sizeof(buf_));
    ::new (static_cast<void*>(buf_)) Deleter(d);
}
```

- Constructor does three things:
 - Store the object pointer p
 - Store the deleter in the buffer
 - Point vtable to the right static variable

Type erasure implementation #3

```
template <typename T> class smartptr_te_vtable {  
    const vtable_t* vtable_ = nullptr;  
    alignas(8) char buf_[8];  
public:  
    ~smartptr_te_vtable() {  
        this->vtable_->destroy_(p_, buf_);  
        this->vtable_->destructor_(buf_);  
    }  
};
```

- Copy etc are handled similarly
- Only one vtable pointer in the class!



operations on
erased type

```
struct vtable_t {  
    using destroy_t = void(*) (T*, void*);  
    using destructor_t = void(*) (void*);  
    destroy_t destroy_;  
    destructor_t destructor_;  
};
```

Type erasure using v-table

- At the point where the type is erased, the compiler generates multiple correct-by-construction reification functions
 - The erased type is hidden in the code of these functions
- The signature of all functions is type-agnostic (`void*`)
- All function pointers are stored in a static variable
 - Template static variable, depends on the deleter, constructor instantiates it
- The vtable pointer is set to the right static vtable variable
- The deleter is saved in a buffer (local or dynamic)
- Type `Deleter` has been erased: we have `f(void*)` and `char[]`
- All reification functions are invoked through their pointers in the vtable
- This really is how compilers build v-tables!

Type erasure using v-tables

- At the point where the type is erased, the compiler instantiates a static variable that depends on the erased type
- Initializing this variable instantiates function template on the same type
 - Erased type is hidden in the generated code
- All function template instantiations have the same signature
- All static variable instantiations have the the same type
- Objects with state are stored in local or dynamic memory
- Type-erased code is executed by an (double) indirect function call
- For each supported operation, a function pointer in the vtable is needed
- There is only one vtable pointer in the object

Outline

- What is type erasure
- What does type erasure look like
- Type erasure as a design pattern
- Type erasure as an implementation technique
- How does it work?
- Three ways to implement type erasure in C++
 - Inheritance
 - Static functions
 - V-table
- Performance benchmarks

Performance of type erasure

- Smart pointer creation and deletion

Benchmark	Time	CPU	Iterations	UserCounters...
BM_rawptr	8.88 ns	8.88 ns	81125063	items_per_second=112.561M/s
BM_uniqueptr	9.12 ns	9.12 ns	76533789	items_per_second=109.678M/s
BM_sharedptr	21.5 ns	21.5 ns	32544564	items_per_second=46.5102M/s
BM_make_sharedptr	12.0 ns	12.0 ns	58151720	items_per_second=83.2405M/s
BM_smartptr	9.37 ns	9.34 ns	76504653	items_per_second=107.093M/s
BM_smartptr_te	19.7 ns	19.7 ns	35678838	items_per_second=50.8338M/s
BM_smartptr_te_lb_opt	10.6 ns	10.6 ns	65890276	items_per_second=94.0315M/s
BM_smartptr_te_lb_only	10.5 ns	10.5 ns	65952774	items_per_second=94.8626M/s
BM_smartptr_te_static	9.86 ns	9.86 ns	70883148	items_per_second=101.38M/s
BM_smartptr_te_vtable	10.6 ns	10.6 ns	60152866	items_per_second=94.6588M/s

Performance of type erasure

- Deleter performance of smart pointers is not all that exciting
- What's the most performance-critical type-erased object?
- `std::function`
 - Type erasure machinery is exercised on every call!
- How to implement a type-erased function?

Performance of type erasure

- Deleter performance of smart pointers is not all that exciting
- What's the most performance-critical type-erased object?
- `std::function`
 - Type erasure machinery is exercised on every call!
- How to implement a type-erased function?

Performance of type erasure

- How to implement a type-erased function?
- Use the fastest option 2 (static function) for the call
- Use option 3 (vtable) for copy, move, delete, etc (no object to store)
- Use local buffer optimization (only? - design decision)
- Optimize for trivially-everything objects? (another design decision)

Implementation based on works by Arthur O'Dwyer and Eduardo Madrid

- Zoo project (<https://github.com/thecppzoo/zoo>)

Type-erased `nonstd::function`

```
template<typename Signature, size_t Size = 16,  
        size_t Alignment = 8> struct Function;
```

- Size and Alignment are for the local buffer (ignore for SlideWare)

```
template<typename Signature> class Function;
```

incomplete type

```
template<typename Res, typename... Args>
```

```
struct Function<Res(Args...)> {...};
```

- Partial specialization is convenient to extract return type and parameters
 - Reflection, yay..!

Type-erased `nonstd::function`

- Local buffer (remember Size and Alignment?)

```
template<typename Res, typename... Args>
struct Function<Res(Args...)> {
    alignas(Alignment) char space_[Size];
    ...
};
```

- This is where we store the callable object
 - `std::function` has an 8-byte buffer
 - enough for function pointers and stateless callables like lambdas
 - function pointers require 16 bytes

Type-erased nonstd::function

- Function call uses the static function method:


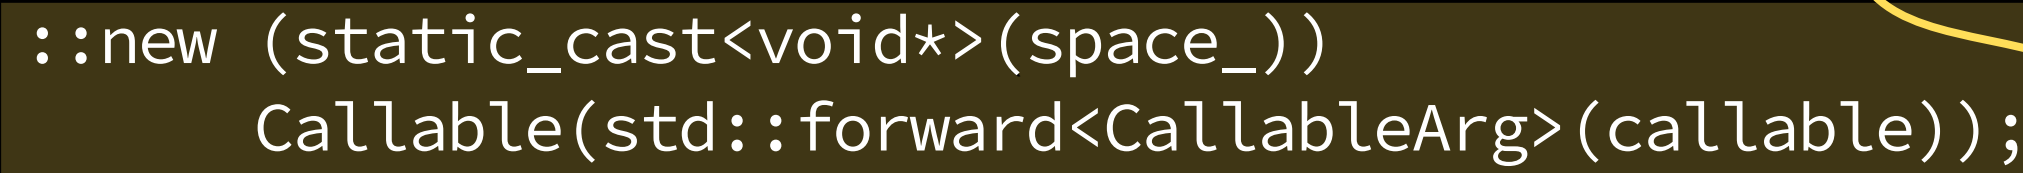
```
using executor_t = Res(*) (Args..., Function*);  
executor_t executor_;  
template<typename Callable>  
static Res executor(Args... args, Function* this_function)  
{  
    return (*reinterpret_cast<Callable*>(this_function->space_))  
           (std::forward<Args>(args)...);  
}
```



Annotations:

- type-agnostic signature
- function pointer
- reification template function
- restore Callable
- invoke Callable with args

Type-erased nonstd::function

- Constructor hides erased type in the code it generates:


```
template <typename CallableArg,  
         typename Callable = std::decay_t<CallableArg>>  
    requires(!std::same_as<Function, Callable>)   
    Function(CallableArg&& callable) : executor_(executor<Callable>)  
    {  
          
        Callable(std::forward<CallableArg>(callable));  
    }
```

- Store the callable in the buffer (strip references) 
- Generate reification function and set the function pointer (executor_) 

Type-erased nonstd::function

- The call operator invokes the executor with the specified arguments
- The arguments do not have to match the function signature but must be convertible to those
 - use concepts or static asserts for better error messages

```
template <typename... CallArgs>
Res operator()(CallArgs&&... callargs) const {
    return this->executor_(std::forward<CallArgs>(callargs)...,
                           const_cast<Function*>(this));
}
```



template instantiation
created during construction

Type-erased nonstd::function

- How to deal with copy, move, and destruction?
 - 1) Implement using vtable
 - 2) Restrict to trivially-everything types

```
template <typename CallableArg,  
         typename Callable = std::decay_t<CallableArg>>  
Function(CallableArg&& callable) : executor_(executor<Callable>) {  
    static_assert(sizeof(Callable) <= Size);  
    static_assert(alignof(Callable) <= Alignment);  
    static_assert(std::is_trivially_destructible<Callable>::value);  
    static_assert(std::is_trivially_copyable<Callable>::value);  
    ...  
}
```

Type-erased `nonstd::function`

- `std::function` can be default-constructed (nothing to call)
 - Throws `std::bad_function_call` if called anyway
- Bad: default `executor_` to null and check at run-time
 - Check is done for all functions, initialized or not
- Good: `executor_` is never null, default executor throws

```
static constexpr Res default_executor(Args..., Function*) {  
    throw std::bad_function_call();  
}  
  
constexpr static executor_t default_executor=default_executor;  
executor_t executor_ = default_executor;
```

Type-erased `nonstd::function`

- Destruction, copying, etc are handled by the vtable
- Member functions can be trivially supported
 - needs a constructor overload and another executor template
- Dynamic buffers for large callables are straightforward
- None of these affect performance of the function call
 - Local buffer might, so compare fairly (`std::function` also uses buffer)

Type erasure using static functions (again)

- At the point where the type is erased, the compiler instantiates a static function template that depends on the erased type
 - Erased type is hidden in the generated code
- All template instantiations have the same signature
- Template instantiation is assigned to a function pointer
- Default function pointer assignment performs the default action
- Objects with state are stored in local or dynamic memory
- Type-erased code is executed by an indirect function call
- Other, less performance-critical operations are handled using vtable

Type-erased nonstd::function performance

- Let's see what a call to a function looks like:

```
int f(int a, int b, int c, int d);  
using F = int(int, int, int, int);  
auto F_invoke(int a, int b, int c, int d, F f) {  
    return f(a, b, c, d);  
}
```

- Assembly of F_invoke:

```
000000000000000000 <_Z8F_invokeiiiiPFiiiiE>:  
    0:  41 ff e0                jmpq    *%r8
```

Type-erased nonstd::function performance

- Now let's see what a call to a std::function looks like:

```
int f(int a, int b, int c, int d);
using F = int(int, int, int, int);
using SF = std::function<F>;
auto SF_invoke(int a, int b, int c, int d, const SF& f) {
    return f(a, b, c, d);
}
```

- Assembly of SF_invoke:

0000000000000000 <_ZSF_invoke>:

0:	48 83 ec 18	sub \$0x18,%rsp
4:	49 83 78 10 00	cmpq \$0x0,0x10(%r8)
9:	89 3c 24	mov %edi,(%rsp)
c:	89 74 24 04	mov %esi,0x4(%rsp)
10:	89 54 24 08	mov %edx,0x8(%rsp)
14:	89 4c 24 0c	mov %ecx,0xc(%rsp)
18:	74 20	je 3a<_ZSF_invoke>
1a:	4c 89 c0	mov %r8,%rax

1d:	48 8d 4c 24 08	lea 0x8(%rsp),%rcx
22:	4c 8d 44 24 0c	lea 0xc(%rsp),%r8
27:	48 89 e6	mov %rsp,%rsi
2a:	48 8d 54 24 04	lea 0x4(%rsp),%rdx
2f:	48 89 c7	mov %rax,%rdi
32:	ff 50 18	callq *0x18(%rax)
35:	48 83 c4 18	add \$0x18,%rsp
39:	c3	retq
3a:	e8 00 00 00 00	callq 3f<_ZSF_invoke>

Type-erased nonstd::function performance

- Assembly of SF_invoke:

0000000000000000 <_Z9SF_invoke>:

0:	48 83 ec 18	sub	\$0x18,%rsp	1d:	48 8d 4c 24 08	lea	0x8(%rsp),%rcx
4:	49 83 78 10 00	cmpq	\$0x0,0x10(%r8)	22:	4c 8d 44 24 0c	lea	0xc(%rsp),%r8
9:	89 3c 24	mov	%edi, (%rsp)	27:	48 89 e6	mov	%rsp,%rsi
c:	89 74 24 04	mov	%esi,0x4(%rsp)	2a:	48 8d 54 24 04	lea	0x4(%rsp),%rdx
10:	89 54 24 08	mov	%edx,0x8(%rsp)	2f:	48 89 c7	mov	%rax,%rdi
14:	89 4c 24 0c	mov	%ecx,0xc(%rsp)	32:	ff 50 18	callq	*0x18(%rax)
18:	74 20	je	3a <_Z9SF_invoke>	35:	48 83 c4 18	add	\$0x18,%rsp
1a:	4c 89 c0	mov	%r8,%rax	39:	c3	retq	
				3a:	e8 00 00 00 00	callq	3f <_Z9SF_invoke>

Type-erased nonstd::function performance

- OK, so what does a call to our type-erased function looks like?

```
int f(int a, int b, int c, int d);  
using F = int(int, int, int, int);  
using FF = Function<F>;  
auto FF_invoke(int a, int b, int c, int d, const FF& f) {  
    return f(a, b, c, d);  
}
```

- Assembly of FF_invoke:

```
00000000000000000000 <_Z9FF_invokeiiiiRK8FunctionEE>:  
    0:  41 ff 60 10          jmpq    *0x10(%r8)
```


Type-erased `nonstd::function` real performance

- Better assembly does not always translate into better performance
- We must benchmark the call itself

Benchmark	Time	CPU	Iterations	UserCounters...
<hr/>				
BM_F_invoke	25.3 ns	25.3 ns	27518802	items_per_second=1.26442G/s
BM_FF_invoke	26.0 ns	26.0 ns	26985686	items_per_second=1.22865G/s
BM_SF_invoke	53.7 ns	53.7 ns	12798869	items_per_second=596.021M/s

Type-erased `nonstd::function` real performance

- How does it compare with a regular or virtual function call?
 - The function body is in a separate compilation unit in all cases

Benchmark	Time	CPU	Iterations	UserCounters...
BM_F_invoke	25.3 ns	25.3 ns	27518802	items_per_second=1.26442G/s
BM_FF_invoke	26.0 ns	26.0 ns	26985686	items_per_second=1.22865G/s
BM_SF_invoke	53.7 ns	53.7 ns	12798869	items_per_second=596.021M/s
BM_f	26.0 ns	26.0 ns	27341020	items_per_second=1.23192G/s
BM_virtual	25.9 ns	25.9 ns	27123930	items_per_second=1.236G/s

Type-erased `nonstd::function` real performance

- The cost of a (well-done) indirection is about the same
- Nothing beats the performance boost from inlining

Benchmark	Time	CPU	Iterations	UserCounters...
BM_F_invoke	25.3 ns	25.3 ns	27518802	items_per_second=1.26442G/s
BM_FF_invoke	26.0 ns	26.0 ns	26985686	items_per_second=1.22865G/s
BM_SF_invoke	53.7 ns	53.7 ns	12798869	items_per_second=596.021M/s
BM_f	26.0 ns	26.0 ns	27341020	items_per_second=1.23192G/s
BM_inline	0.209 ns	0.209 ns	1000000000	items_per_second=153.049G/s
BM_virtual	25.9 ns	25.9 ns	27123930	items_per_second=1.236G/s

How is type erasure done in C++?

- Type erasure in C++ is very similar to C:
- The generic code does not have any mention of the erased type
 - Often uses `void*` or `char*`
- The erased type is hidden in the code of a function that is invoked to perform the type-dependent action
 - The signature of this function is type-agnostic
 - The body of this function reifies the erased type (often with casts)
- C++ automates writing the reification code and ensures that it matches the erased type
- The code with the hidden type is generated by a template at the point where the erased type is last present

What is type erasure?

- Type erasure is used to separate the interface from the implementation
 - Even more: separate relevant interface (type properties) from the rest
 - Other than having the relevant interface, types can be very different
- Type erasure can be used to implement separation of concerns
- Type erasure is often used to break dependencies
- Type erasure doesn't have to be any more expensive than any other indirection mechanism
 - [with a good implementation] there is no overhead assuming the indirection was needed
- Indirection can be expensive in any guise
- Is decoupling worth the cost of indirection? That is a design decision

C++ Type Erasure Demystified

Questions?

Possibly answers too...