

C++ now

# Security in C++

## *Hardening Techniques from the Trenches*

Louis Dionne

2024

# Who Am I Quoting?

Experts have identified a few programming languages that both **lack traits associated with memory safety** and also have high proliferation across critical systems, such as **C and C++**.

Choosing to use memory safe programming languages at the outset, as recommended by the Cybersecurity and Infrastructure Security Agency's (CISA) Open-Source Software Security Roadmap is one example of developing software in a secure-by-design manner. \*

# How Important Is Security?

- Cybercrime is estimated to cost 10 trillion by 2025
- Spyware like Pegasus targeted journalists and high-profile activists
- The WannaCry ransomware affected over 300k computers in 150 countries, cost over \$4B
  - Affected critical systems like hospitals

# C++ Is Memory Unsafe

**Memory Unsafety Accounts  
For ~70% Of High Severity  
Security Issues \***

# We Can Do Better

# We Must Do Better

# Partly an Attitude Problem

C++ has generally adopted an expert-friendly attitude:

- If the user makes a mistake, it's their fault
- Performance at all costs
- Complicated APIs in the name of flexibility

# The Mindset Is Changing

- More general awareness about the problem
- Creation of SG23 (Safety and Security Study Group)
- Most of WG21 understands this is a make or break
- However, still few concrete improvements

We're Engineers, Let's Solve  
Problems

# Agenda

## Overview of Memory Safety

Library Undefined Behavior

Standard Library Hardening

Typed Memory Operations

Conclusions

# Types of Memory Safety

- **Spatial memory safety**
- **Temporal memory safety**
- Type safety
- Guaranteed initialization
- Thread safety

# Spatial memory safety

- Each memory allocation has a given size (or bounds)
- Accessing memory out of bounds is called an out-of-bounds (OOB) access

# Example

```
int main() {
    char input[8];
    char password[8];

    std::ifstream("/etc/password") >> password;
    std::cout << "Enter password: ";
    std::cin >> input;

    if (std::strcmp(password, input, 8) == 0)
        std::cout << "Access granted";
    else
        std::cout << "Access denied";
}
```

```
$ authenticate
> securitysecurity
```

input	password

# Example

```
int main() {
    char input[8];
    char password[8];

    std::ifstream("/etc/password") >> password;

    std::cout << "Enter password: ";
    std::cin >> input;

    if (std::strcmp(password, input, 8) == 0)
        std::cout << "Access granted";
    else
        std::cout << "Access denied";
}
```

```
$ authenticate
> securitysecurity
```

input	password
	r e a l p a s s

# Example

```
int main() {
    char input[8];
    char password[8];

    std::ifstream("/etc/password") >> password;
    std::cout << "Enter password: ";
    std::cin >> input;

    if (std::strcmp(password, input, 8) == 0)
        std::cout << "Access granted";
    else
        std::cout << "Access denied";
}
```

```
$ authenticate
> securitysecurity
```

input	password
s e c u r i t y r e a l	p a s s

# Example

```
int main() {
    char input[8];
    char password[8];

    std::ifstream("/etc/password") >> password;
    std::cout << "Enter password: ";
    std::cin >> input;

    if (std::strcmp(password, input, 8) == 0)
        std::cout << "Access granted";
    else
        std::cout << "Access denied";
}
```

```
$ authenticate
> security|security
```

input	password
s e c u r i t y	s e c u r i t y



# Example

```
int main() {
    char input[8];
    char password[8];

    std::ifstream("/etc/password") >> password;
    std::cout << "Enter password: ";
    std::cin >> input;

    if (std::strcmp(password, input, 8) == 0)
        std::cout << "Access granted";
    else
        std::cout << "Access denied";
}
```



input	password
s e c u r i t y	s e c u r i t y



```
$ authenticate
> security|security
```

# Now Imagine This

```
using Authentifier = bool (*)(Account const&, Password const&);

struct Account {
    ...
    int transit_number[5];
    int account_number[8];
    Authentifier identification_method = &two_factor_auth;
    ...
};

void teller(FILE* socket) {
    while (true) {
        char const* data = read_socket(socket); // untrusted

        Account account;
        parse_transit_number(data, &account.transit_number);
        parse_account_number(data, &account.account_number);
        auto authenticate = account.identification_method;
        if (authenticate(account, password)) {
            // we're in!
        }
        free(data);
    }
}
```

# Now Imagine This

```
using Authentifier = bool (*)(Account const&, Password const&);

struct Account {
    ...
    int transit_number[5];
    int account_number[8];
    Authentifier identification_method = &two_factor_auth;
    ...
};

void teller(FILE* socket) {
    while (true) {
        char const* data = read_socket(socket); // untrusted

        Account account;
        parse_transit_number(data, &account.transit_number);
        parse_account_number(data, &account.account_number);
        auto authenticate = account.identification_method;
        if (authenticate(account, password)) {
            // we're in!
        }
        free(data);
    }
}
```

# Now Imagine This

```
using Authentifier = bool (*)(Account const&, Password const&);

struct Account {
    ...
    int transit_number[5];
    int account_number[8];
    Authentifier identification_method = &two_factor_auth;
    ...
};

void teller(FILE* socket) {
    while (true) {
        char const* data = read_socket(socket); // untrusted

        Account account;
        parse_transit_number(data, &account.transit_number);
        parse_account_number(data, &account.account_number);
        auto authenticate = account.identification_method;
        if (authenticate(account, password)) {
            // we're in!
        }
        free(data);
    }
}
```

# Now Imagine This

```
using Authentifier = bool (*)(Account const&, Password const&);

struct Account {
    ...
    int transit_number[5];
    int account_number[8];
    Authentifier identification_method = &two_factor_auth;
    ...
};

void teller(FILE* socket) {
    while (true) {
        char const* data = read_socket(socket); // untrusted

        Account account;
        parse_transit_number(data, &account.transit_number);
        parse_account_number(data, &account.account_number);
        auto authenticate = account.identification_method;
        if (authenticate(account, password)) {
            // we're in!
        }
        free(data);
    }
}
```

# Now Imagine This

```
using Authentifier = bool (*)(Account const&, Password const&);

struct Account {
    ...
    int transit_number[5];
    int account_number[8];
    Authentifier identification_method = &two_factor_auth;
    ...
};

void teller(FILE* socket) {
    while (true) {
        char const* data = read_socket(socket); // untrusted

        Account account;
        parse_transit_number(data, &account.transit_number);
        parse_account_number(data, &account.account_number);
        auto authenticate = account.identification_method;
        if (authenticate(account, password)) {
            // we're in!
        }
        free(data);
    }
}
```

# Now Imagine This

```
using Authentifier = bool (*)(Account const&, Password const&);

struct Account {
    ...
    int transit_number[5];
    int account_number[8];
    Authentifier identification_method = &two_factor_auth;
    ...
};

void teller(FILE* socket) {
    while (true) {
        char const* data = read_socket(socket); // untrusted

        Account account;
        parse_transit_number(data, &account.transit_number);
        parse_account_number(data, &account.account_number);
        auto authenticate = account.identification_method;
        if (authenticate(account, password)) {
            // we're in!
        }
        free(data);
    }
}
```

# Temporal Memory Safety

- All memory accesses to an object should occur during the lifetime of the object's allocation
- Access to the object outside of this window is called a use-after-free

# Example

```
void teller(FILE* socket) {
    std::vector<std::future<void>> transactions;
    while (true) {
        char const* data = read_socket(socket);
        transactions.push_back(std::async([=] {
            Account account;
            parse_transit_number(data, &account.transit_number);
            parse_account_number(data, &account.account_number);
            auto authenticate = account.identification_method;
            if (authenticate(account, password)) {
                // we're in!
            }
        }));
        free(data);
    }
}
```

# Example

```
void teller(FILE* socket) {
    std::vector<std::future<void>> transactions;
    while (true) {
        char const* data = read_socket(socket);
        transactions.push_back(std::async([=] {
            Account account;
            parse_transit_number(data, &account.transit_number);
            parse_account_number(data, &account.account_number);
            auto authenticate = account.identification_method;
            if (authenticate(account, password)) {
                // we're in!
            }
        }));
        free(data);
    }
}
```

# Example

```
void teller(FILE* socket) {
    std::vector<std::future<void>> transactions;
    while (true) {
        char const* data = read_socket(socket);
        transactions.push_back(std::async([=] {
            Account account;
            parse_transit_number(data, &account.transit_number);
            parse_account_number(data, &account.account_number);
            auto authenticate = account.identification_method;
            if (authenticate(account, password)) {
                // we're in!
            }
        }));
        free(data);
    }
}
```

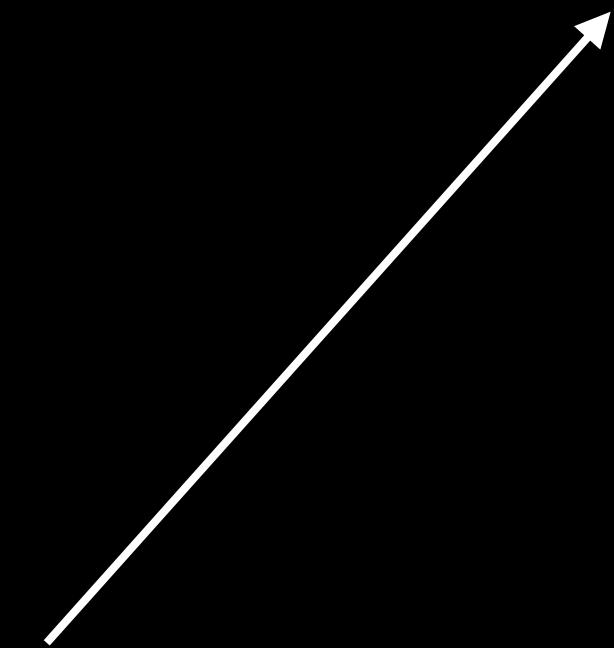
# Example

```
void teller(FILE* socket) {
    std::vector<std::future<void>> transactions;
    while (true) {
        char const* data = read_socket(socket);
        transactions.push_back(std::async([=] {
            Account account;
            parse_transit_number(data, &account.transit_number);
            parse_account_number(data, &account.account_number);
            auto authenticate = account.identification_method;
            if (authenticate(account, password)) {
                // we're in!
            }
        }));
        free(data);
    }
}
```

# Example

```
void teller(FILE* socket) {
    std::vector<std::future<void>> transactions;
    while (true) {
        char const* data = read_socket(socket);
        transactions.push_back(std::async([=] {
            Account account;
            parse_transit_number(data, &account.transit_number);
            parse_account_number(data, &account.account_number);
            auto authenticate = account.identification_method;
            if (authenticate(account, password)) {
                // we're in!
            }
        }));
        free(data);
    }
}
```

read\_socket's  
allocation

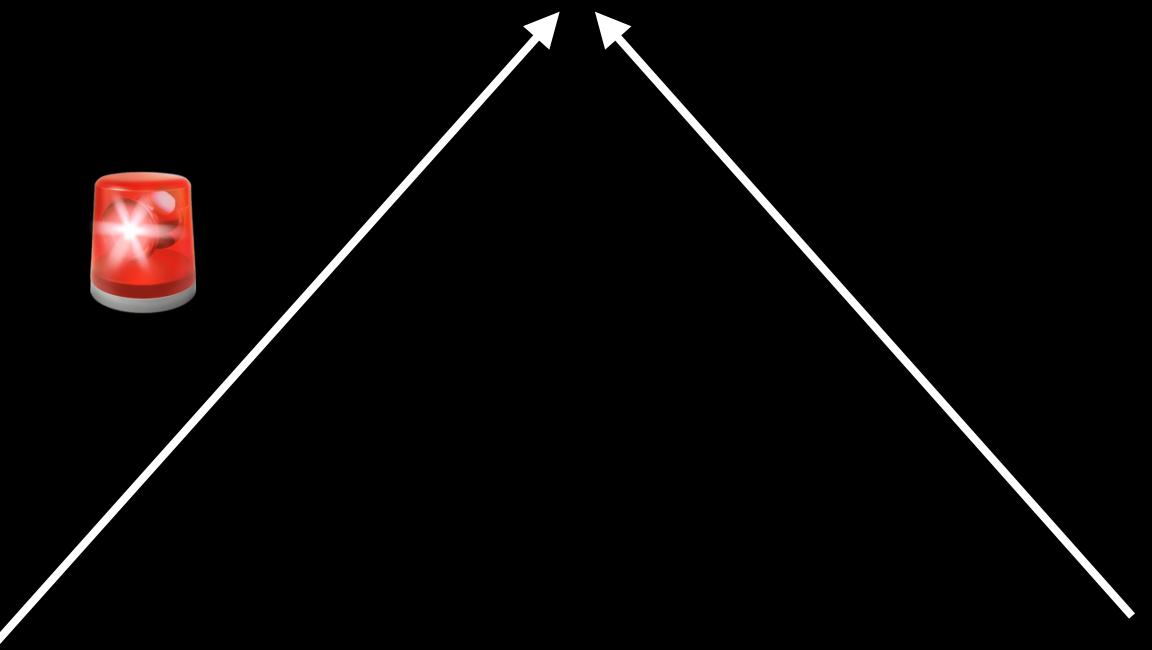


Transaction 1

# Example

```
void teller(FILE* socket) {
    std::vector<std::future<void>> transactions;
    while (true) {
        char const* data = read_socket(socket);
        transactions.push_back(std::async([=] {
            Account account;
            parse_transit_number(data, &account.transit_number);
            parse_account_number(data, &account.account_number);
            auto authenticate = account.identification_method;
            if (authenticate(account, password)) {
                // we're in!
            });
        }));
        free(data);
    }
}
```

read\_socket's  
allocation



Transaction 1

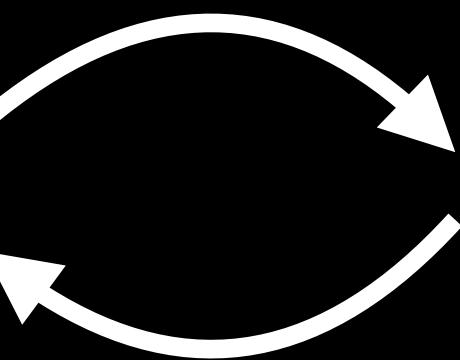
Transaction 2

# Type Safety

- A memory allocation is used to represent an object of a particular type
- Interpreting it as an object of a different type is called a type confusion

# Most Temporal Memory Issues Involve “Type Confusion”

```
struct timespec {  
    time_t tv_sec;  
    long tv_nsec;  
};
```



```
struct iovec {  
    char* iov_base;  
    size_t iov_len;  
};
```

# Guaranteed Initialization

- When memory is allocated, it contains garbage
- Using that undefined content can lead to information disclosure
- Can also be exploited if the attacker controls the “garbage”

# Thread Safety

- Multi-threaded programs do concurrent accesses to memory
- If accesses are not properly synchronized, they can lead to data races
- Those can be composed (typically non-deterministically) with other bugs to create exploits
- Techniques exist to make non-deterministic exploits fairly reliable

# What About C++?

It turns out that most safety issues are technically Undefined Behavior

# Agenda

Overview of Memory Safety

**Library Undefined Behavior**

Standard Library Hardening

Typed Memory Operations

Conclusions

# Library Undefined Behavior

UB is what happens when the Standard doesn't guarantee anything  
For the library, it is generally even stated explicitly:

24 Containers library		[containers]
24.7	Views	[views]
24.7.2	Contiguous access	[views.contiguous]
24.7.2.2	Class template <code>span</code>	[views.span]
24.7.2.6	Element access	[span.elem]
<pre>constexpr reference operator[](size_type idx) const;</pre>		
1	<i>Preconditions:</i> <code>idx &lt; size()</code> is true.	←
2	<i>Effects:</i> Equivalent to: <code>return *data() + idx;</code>	
<pre>constexpr reference front() const;</pre>		
3	<i>Preconditions:</i> <code>empty()</code> is false.	←
4	<i>Effects:</i> Equivalent to: <code>return *data();</code>	
<pre>constexpr reference back() const;</pre>		
5	<i>Preconditions:</i> <code>empty()</code> is false.	←
6	<i>Effects:</i> Equivalent to: <code>return *(data() + (size() - 1));</code>	
<pre>constexpr pointer data() const noexcept;</pre>		
7	<i>Effects:</i> Equivalent to: <code>return data_;</code>	

# Undefined Behavior Is a Specification Tool

- Creates a contract with the programmer
- Allows writing simpler APIs that make sense  
(No need to define results for meaningless inputs)
- Gives freedom to the implementation

# Implementations Can Take Advantage of It

Valid strategies:

- Don't do anything (execute code by assuming preconditions)
- Trap if a precondition is violated
- Log a message, keep executing
- ... basically anything else you can think of

# Agenda

Overview of Memory Safety

Library Undefined Behavior

**Standard Library Hardening**

Typed Memory Operations

Conclusions

# Standard Library Hardening

- Turn select UB into guaranteed traps
- Provide *hardening* modes with high-level semantics
- Allow users to select hardening mode that's right for them
- Allow vendors to select the default mode

This is not a “debugging”  
feature

**You should ship that way!**

# Libc++ Hardening Modes

- **none** — do not compromise *any* performance
- **fast** — security-critical low-overhead checks only
- **extensive** — low-overhead checks
- **debug** — all checks

# Inside the Library

- Checks are grouped in a few large categories
- Categories represent the nature of a check

# Valid Element Access Checks

Checks that an attempt to access a container element is valid

`std::optional` is considered a container

Example:

```
template <class _Tp, class _Allocator>
const_reference vector<_Tp, _Allocator>::operator[](size_type __n) const noexcept {
    _LIBCPP_ASSERT_VALID_ELEMENT_ACCESS(__n < size(), "vector[] index out of bounds");
    return this->__begin_[__n];
}
```

# Valid Element Access Checks

Checks that an attempt to access a container element is valid

`std::optional` is considered a container

Example:

```
template <class _Tp, class _Allocator>
const_reference vector<_Tp, _Allocator>::operator[](size_type __n) const noexcept {
    _LIBCPP_ASSERT_VALID_ELEMENT_ACCESS(__n < size(), "vector[] index out of bounds");
    return this->__begin_[__n];
}
```

# Valid Input Range Checks

Checks that a range given as an input is valid:

- End iterator is reachable from the begin iterator
- Both iterators refer to the same container

Example:

```
template <class _CharT, class _Traits, class _Allocator>
iterator basic_string<_CharT, _Traits, _Allocator>::erase(const_iterator __first, const_iterator __last) {
    _LIBCPP_ASSERT_VALID_INPUT_RANGE(__first <= __last, "string::erase(first, last) called with invalid range");
    iterator __b = begin();
    size_type __r = static_cast<size_type>(__first - __b);
    erase(__r, static_cast<size_type>(__last - __first));
    return __b + static_cast<difference_type>(__r);
}
```

# Valid Input Range Checks

Checks that a range given as an input is valid:

- End iterator is reachable from the begin iterator
- Both iterators refer to the same container

Example:

```
template <class _CharT, class _Traits, class _Allocator>
iterator basic_string<_CharT, _Traits, _Allocator>::erase(const_iterator __first, const_iterator __last) {
    _LIBCPP_ASSERT_VALID_INPUT_RANGE(__first <= __last, "string::erase(first, last) called with invalid range");
    iterator __b = begin();
    size_type __r = static_cast<size_type>(__first - __b);
    erase(__r, static_cast<size_type>(__last - __first));
    return __b + static_cast<difference_type>(__r);
}
```

# Internal Checks

Internal preconditions that do not relate to user code

A failure here is a bug in libc++, not a bug in the user code

Example:

```
template <class _NodePtr>
_NodePtr __tree_min(_NodePtr __x) noexcept {
    _LIBCPP_ASSERT_INTERNAL(__x != nullptr, "Root node shouldn't be null");
    while (__x->__left_ != nullptr)
        __x = __x->__left_;
    return __x;
}
```

# Internal Checks

Internal preconditions that do not relate to user code

A failure here is a bug in libc++, not a bug in the user code

Example:

```
template <class _NodePtr>
_NodePtr __tree_min(_NodePtr __x) noexcept {
    _LIBCPP_ASSERT_INTERNAL(__x != nullptr, "Root node shouldn't be null");
    while (__x->__left_ != nullptr)
        __x = __x->__left_;
    return __x;
}
```

# High Level Modes are Collections of Categories

valid-input-range			
valid-element-access			
non-null-argument			
non-overlapping-ranges			
compatible-allocator			
valid-comparator			
internal			

# High Level Modes are Collections of Categories

	None			
valid-input-range	✗			
valid-element-access	✗			
non-null-argument	✗			
non-overlapping-ranges	✗			
compatible-allocator	✗			
valid-comparator	✗			
internal	✗			

# High Level Modes are Collections of Categories

	None	Fast
valid-input-range	✗	✓
valid-element-access	✗	✓
non-null-argument	✗	✗
non-overlapping-ranges	✗	✗
compatible-allocator	✗	✗
valid-comparator	✗	✗
internal	✗	✗

# High Level Modes are Collections of Categories

	None	Fast	Extensive
valid-input-range	✗	✓	✓
valid-element-access	✗	✓	✓
non-null-argument	✗	✗	✓
non-overlapping-ranges	✗	✗	✓
compatible-allocator	✗	✗	✓
valid-comparator	✗	✗	✗
internal	✗	✗	✗

# High Level Modes are Collections of Categories

	None	Fast	Extensive	Debug
valid-input-range	✗	✓	✓	✓
valid-element-access	✗	✓	✓	✓
non-null-argument	✗	✗	✓	✓
non-overlapping-ranges	✗	✗	✓	✓
compatible-allocator	✗	✗	✓	✓
valid-comparator	✗	✗	✗	✓
internal	✗	✗	✗	✓

# Selecting the Hardening Mode

- Define this macro: `-D_LIBCPP_HARDENING_MODE=<mode>`
- Valid modes are:
  - `_LIBCPP_HARDENING_MODE_NONE`
  - `_LIBCPP_HARDENING_MODE_FAST`
  - `_LIBCPP_HARDENING_MODE_EXTENSIVE`
  - `_LIBCPP_HARDENING_MODE_DEBUG`
- Hardening mode can be selected in each TU

# Usage Example

```
int main(int argc, char** argv) {
    std::vector<std::string> args(argv + 1, argv + argc);

    int total = 0;
    for (int i = 0; i != argc; ++i) {
        total += std::atoi(args[i].c_str());
    }
    std::cout << total << std::endl;
}
```

# Usage Example

```
int main(int argc, char** argv) {
    std::vector<std::string> args(argv + 1, argv + argc);

    int total = 0;
    for (int i = 0; i != argc; ++i) {
        total += std::atoi(args[i].c_str());
    }
    std::cout << total << std::endl;
}
```

# Usage Example

```
int main(int argc, char** argv) {
    std::vector<std::string> args(argv + 1, argv + argc);

    int total = 0;
    for (int i = 0; i != argc; ++i) {
        total += std::atoi(args[i].c_str());
    }
    std::cout << total << std::endl;
}
```

# Usage Example

```
int main(int argc, char** argv) {
    std::vector<std::string> args(argv + 1, argv + argc);

    int total = 0;
    for (int i = 0; i != argc; ++i) {
        total += std::atoi(args[i].c_str());
    }
    std::cout << total << std::endl;
}
```

# Usage Example

```
$ clang++ test.cpp -D_LIBCPP_HARDENING_MODE=_LIBCPP_HARDENING_MODE_DEBUG  
$ ./a.out 1 2 3  
[...]/usr/include/c++/v1/vector:1537: assertion __n < size() failed: vector[] index out of bounds  
zsh: abort      ./a.out 1 2 3
```

# Usage Example

```
$ clang++ test.cpp -D_LIBCPP_HARDENING_MODE=_LIBCPP_HARDENING_MODE_FAST  
$ ./a.out 1 2 3  
zsh: trace trap      ./a.out 1 2 3
```

# Verbosity of Handler

Tradeoff between code size/optimizations and user experience

- fast → trap
- extensive → trap
- debug → abort verbosely

# ABI Considerations

Some useful checks require changing the ABI:

```
std::span<int> span(ptr, 3);
auto b = span.begin();
b += 999;
int value = *b; // can we trap here?
```

# ABI Selection Is Orthogonal to Hardening

- ABI is a property of the platform
- Vendors can select the desired ABI
- It doesn't make sense for users to control that
- Huge simplification: this prevents having to deal with ABI-related concerns as part of hardening

# Example: Bounded Iterators

Library is configured with `_LIBCPP_ABI_BOUNDED_ITERATORS` (by vendor)

```
template <class _Tp, size_t _Extent>
class span {
public:
    using element_type          = _Tp;
    using value_type             = remove_cv_t<_Tp>;
    using size_type              = size_t;
    // ...

#ifndef _LIBCPP_ABI_BOUNDED_ITERATORS
    using iterator               = __bounded_iter<pointer>;
#else
    using iterator               = pointer;
#endif
    // ...
};
```

# Example: Bounded Iterators

Library is configured with `_LIBCPP_ABI_BOUNDED_ITERATORS` (by vendor)

```
template <class _Tp, size_t _Extent>
class span {
public:
    using element_type          = _Tp;
    using value_type             = remove_cv_t<_Tp>;
    using size_type              = size_t;
    // ...

#ifndef _LIBCPP_ABI_BOUNDED_ITERATORS
    using iterator               = __bounded_iter<pointer>;
#else
    using iterator               = pointer;
#endif
    // ...
};
```

# Example: Bounded Iterators

Iterators now have enough information for bounds checking:

```
std::span<int> span(ptr, 3);
auto b = span.begin();
b += 999;
int value = *b; // trap!
```

If hardening mode is *none*, there is **still** no trap

# Deployment Experience

Very positive experience so far:

- WebKit is using hardened libc++
- Several other Apple adopters
- Other adoption: Chrome and Google Cloud Network Virtualization \*
- Lots of anecdotal evidence of bugs being found
- Already some known in-the-wild security issues that hardening would have prevented or alleviated

# Deployment Experience

However, adoption can require some work:

- Adoption is easy for modern C++ code bases
- Harder for code bases that don't use the Standard Library
- Adoption of any new feature can introduce bugs if not careful
- Non-zero performance cost

# ⚠ Latent bugs may result in new crashes

## A typical example

```
std::array<T, 30> array = {...};  
std::sort(&array[0], &array[30]);
```

# Diagnostics to Push Adoption

## Clang's -Wunsafe-buffer-usage

```
int& f(int* p, int n) {  
    return p[n];  
}
```

```
$ clang++ foo.cpp -Wunsafe-buffer-usage  
<stdin>:2:10: warning: unsafe buffer access [-Wunsafe-buffer-usage]  
 2 |     return p[n];  
   |           ^  
1 warning generated.
```

# Standardization Path

# Enter Contracts

- Observe that we're only enforcing existing preconditions
- Contracts provide a framework for expressing them

```
template <class _Tp, class _Allocator>
const_reference vector<_Tp, _Allocator>::operator[](size_type __n) const noexcept
    pre [[clang::diag("vector[] index out of bounds")]] (__n < size())
{
    return this->__begin_[__n];
}
```

# P2900 Contracts MVP

P2900 is based on different *contract semantics*:

	verbose	failures trap
ignore	✗	✗
observe	✓	✗
since Tokyo P3191	✗	✓
enforce	✓	✓

# Why We Want Contracts In C++26

- Allows formally expressing library hardening
- Allows formalizing things like sanitizers
  - After all they just check various types of language UB
- Generally useful feature for users to express assertions

# Why We Want Contracts In C++26

A “safe” mode for ISO C++ that safety and security critical industries can build on top of

# Agenda

Overview of Memory Safety

Library Undefined Behavior

Standard Library Hardening

**Typed Memory Operations**

Conclusions

# A Clever Observation

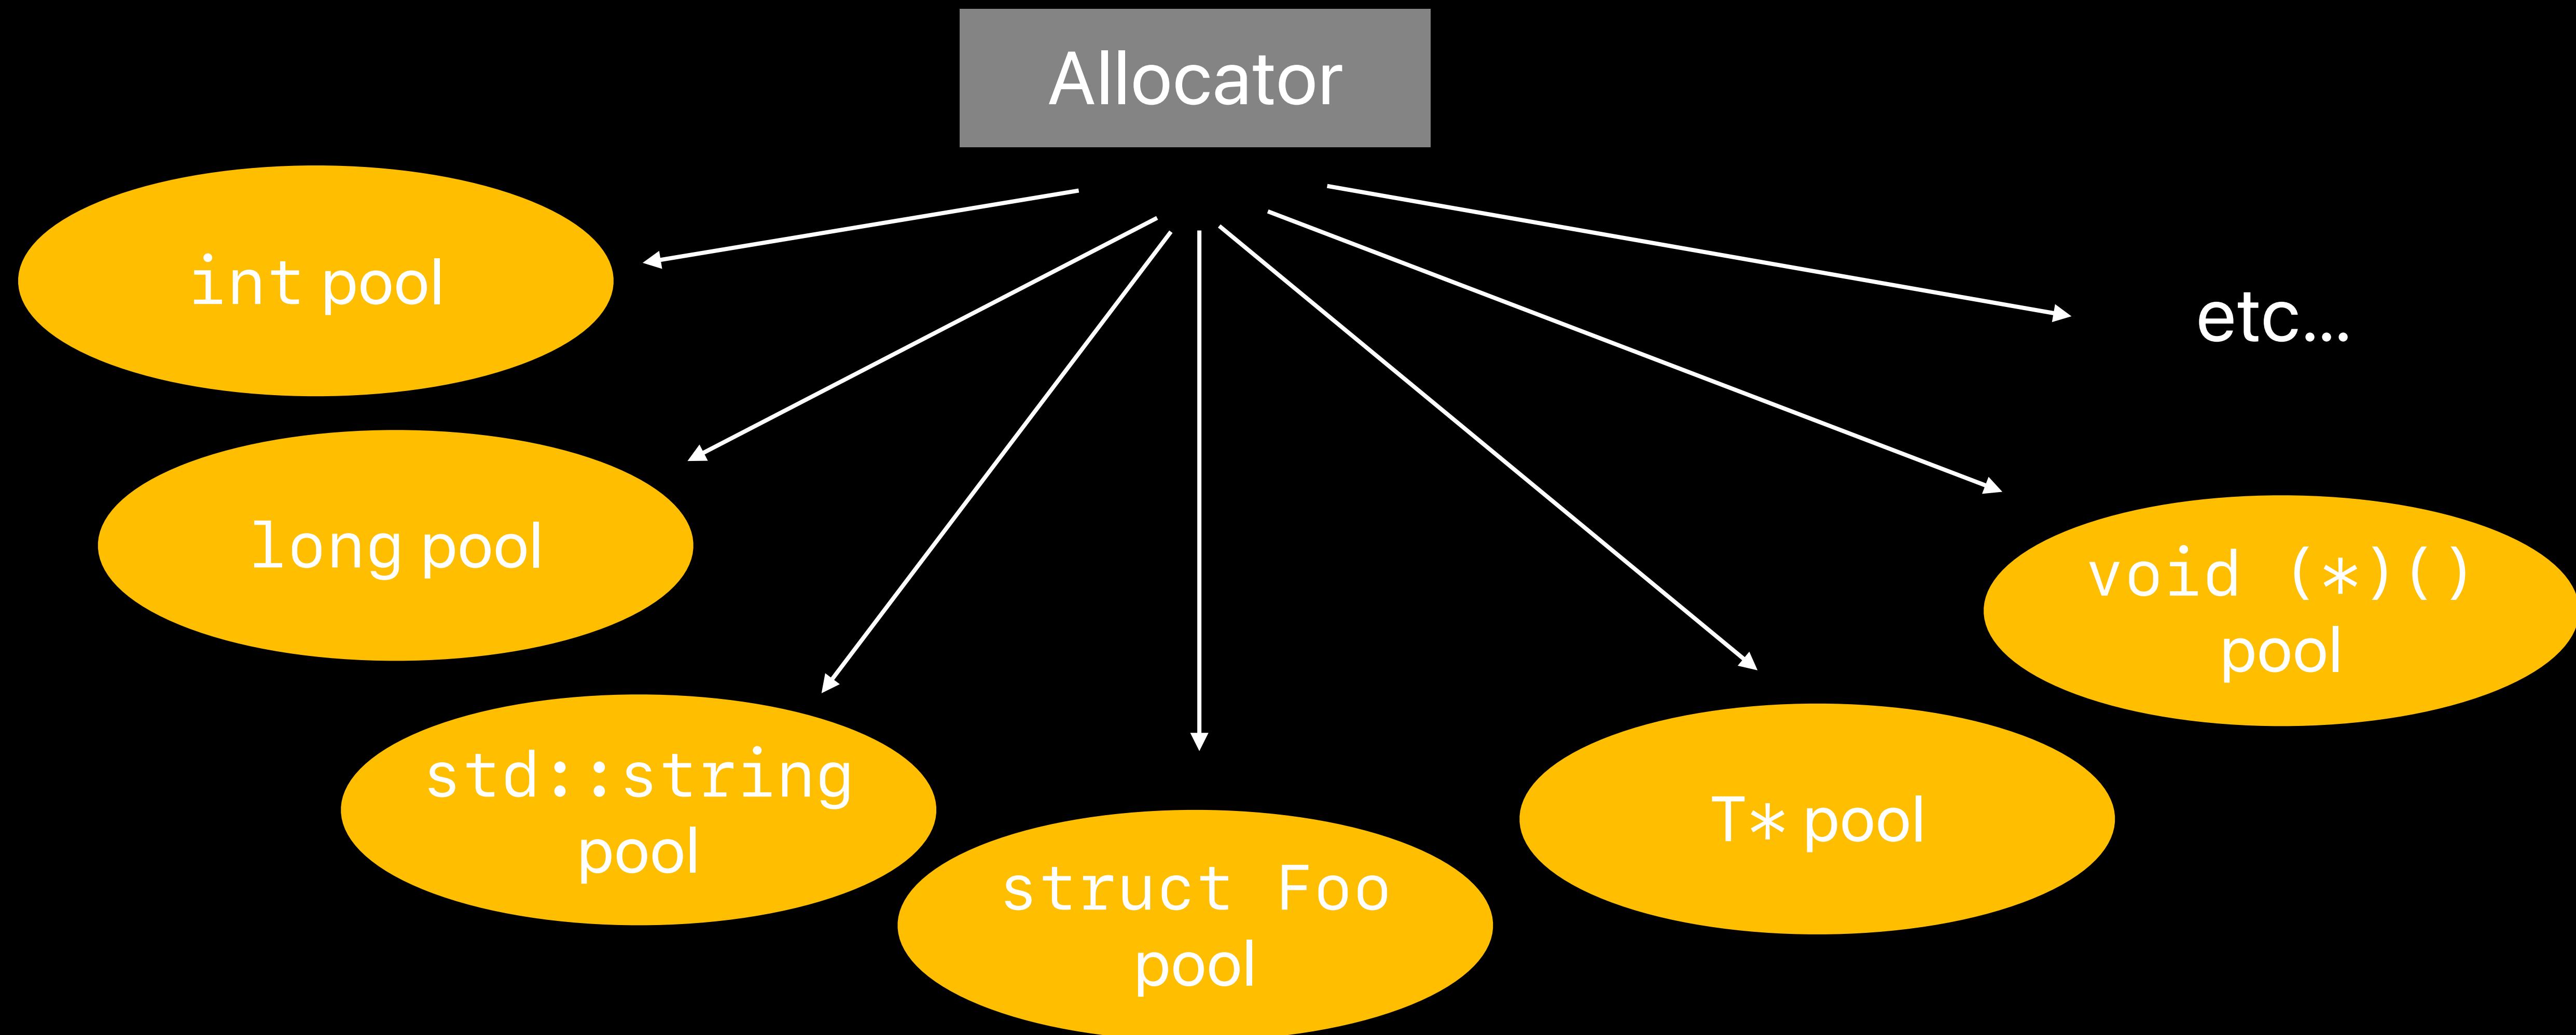
Most temporal memory safety exploits require some type confusion

If memory is never reused for a different type, confusions are impossible

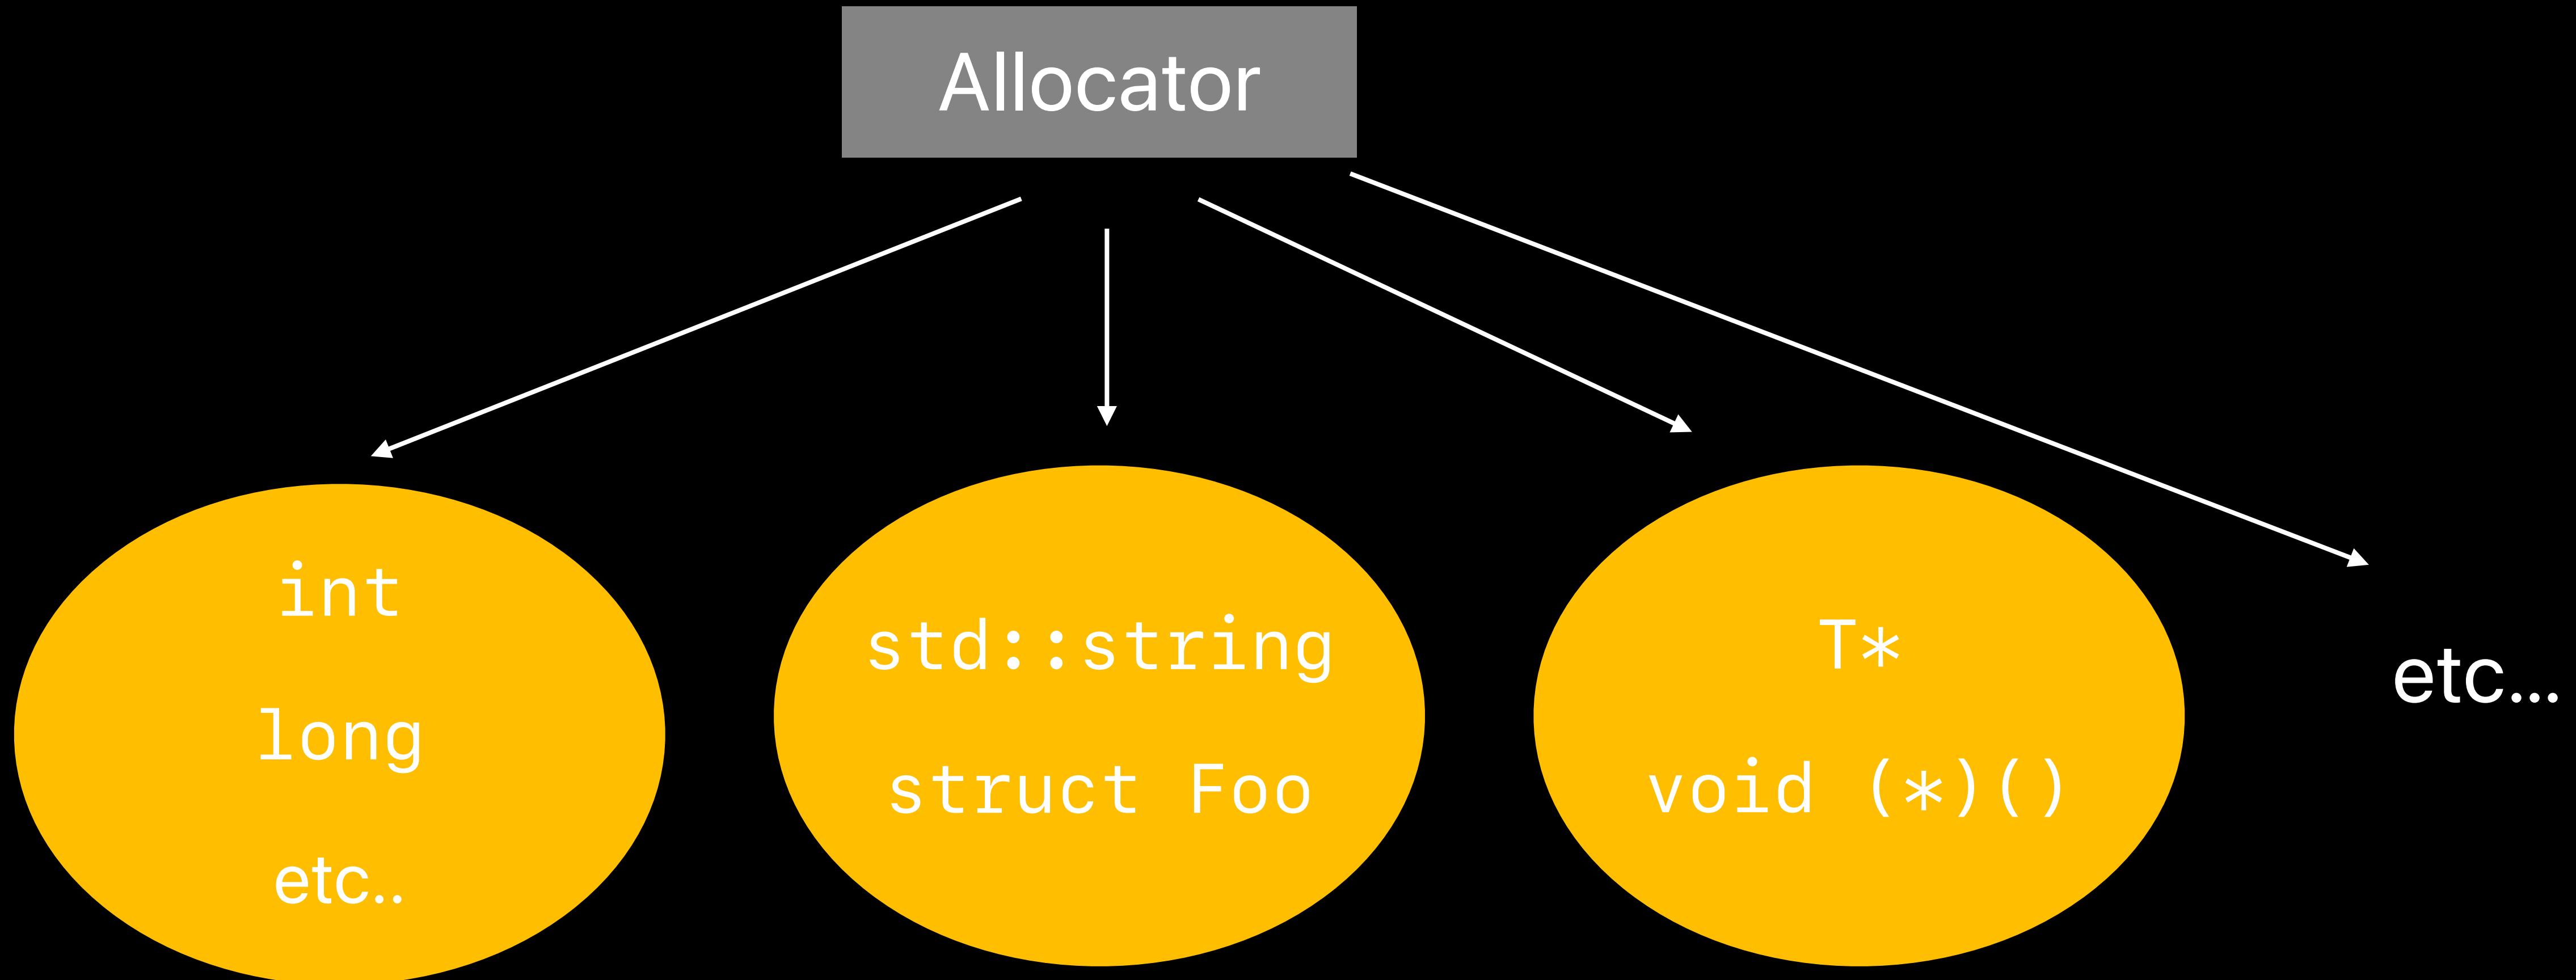
→ Segregate allocations by type!

This was introduced in the Darwin Kernel a few years ago \*

# A Naive Typed Memory Allocator



# A Performance / Security Tradeoff



## The Most Important Property

**Data must not alias pointers**

Making Exploits Harder

**Randomize buckets on boot!**

# How Effective Is Type Segregation?

In the Darwin Kernel, data suggests that the vast majority of dynamic allocation lifetime bugs are not exploitable anymore

# Type Segregation For General C++

```
struct Foo {  
    // ...  
};  
  
std::unique_ptr<Foo> f() {  
    return new Foo{args...}; // GOAL: should come from the Foo pool  
}
```

# The Usual operator new Rewriting

User writes

```
std::unique_ptr<Foo> f() {  
    return new Foo{args...};  
}
```

Compiler rewrites

```
std::unique_ptr<Foo> f() {  
    Foo* __alloc = operator new(sizeof(Foo));  
  
    new (__alloc) Foo{args...};  
    return __alloc;  
}
```

# The Problem

There is no type information

```
void* operator new(std::size_t);
void* operator new(std::size_t, const std::nothrow_t&) noexcept;
void* operator new(std::size_t, std::align_val_t);
void* operator new(std::size_t, std::align_val_t, const std::nothrow_t&) noexcept;
```

# Thankfully, We Can Modify the Standard Library!

Let's add type information

```
enum class __type_descriptor_t : unsigned long long;  
  
void* operator new(std::size_t, std::__type_descriptor_t);  
void* operator new(std::size_t, const std::nothrow_t&, std::__type_descriptor_t) noexcept;
```

# And the Compiler Too!

User writes

```
std::unique_ptr<Foo> f() {  
    return new Foo{args...};  
}
```

Compiler rewrites

```
std::unique_ptr<Foo> f() {  
    Foo* __alloc = operator new(sizeof(Foo));  
  
    new (__alloc) Foo{args...};  
    return __alloc;  
}
```

# And the Compiler Too!

User writes

```
std::unique_ptr<Foo> f() {  
    return new Foo{args...};  
}
```

Compiler rewrites

```
std::unique_ptr<Foo> f() {  
    Foo* __alloc = operator new(sizeof(Foo),  
                               __builtin_type_descriptor(Foo));  
    new (__alloc) Foo{args...};  
    return __alloc;  
}
```

# Then We Forward Type Information to the System Allocator

```
void* operator new(std::size_t size, std::__type_descriptor_t desc) {
    if (size == 0)
        size = 1;

    void* p;
    while ((p = malloc_type_malloc(size, static_cast<malloc_type_id_t>(desc))) == nullptr) {
        // ...
    }
    if (p == nullptr)
        throw std::bad_alloc();
    return p;
}
```

# **What About User-Defined operator new?**

# Isn't That an ABI Break?

# **It Would Be!**

**It Would Be!**

But we're clever

# It Would Be!

But we're clever

```
void* operator new(std::size_t size, std::__type_descriptor_t desc) {
    if (std::__is_function_overridden(static_cast<void* (*)(&operator new)>(&operator new))) {
        return static_cast<void* (*)(&operator new)(size);
    } else {
        // ... the typed implementation ...
    }
}
```

# The Worst Code I've Written?

```
__attribute__((__section__="__TEXT,__lcxx_override")))
void* operator new(size_t __size) noexcept {
    // ...
}
```

```
template <class _Ret, class... _Args>
bool __is_function_overridden(_Ret (*__fptr)(_Args...)) noexcept {
    extern char __segstart __asm("section$start$__TEXT$__lcxx_override");
    extern char __segend __asm("section$end$__TEXT$__lcxx_override");

    uintptr_t __start = (uintptr_t)(&__segstart);
    uintptr_t __end   = (uintptr_t)(&__segend);
    uintptr_t __ptr   = (uintptr_t)(__fptr);

    return __ptr < __start || __ptr > __end;
}
```

# The Worst Code I've Written?

```
attribute((__section__="__TEXT,__lcxx_override"))  
void* operator new(size_t __size) noexcept {  
    // ...  
}
```

```
template <class _Ret, class... _Args>  
bool __is_function_overridden(_Ret (*__fptr)(_Args...)) noexcept {  
    extern char __segstart __asm("section$start$__TEXT$__lcxx_override");  
    extern char __segend __asm("section$end$__TEXT$__lcxx_override");  
  
    uintptr_t __start = (uintptr_t)(&__segstart);  
    uintptr_t __end = (uintptr_t)(&__segend);  
    uintptr_t __ptr = (uintptr_t)(__fptr);  
  
    return __ptr < __start || __ptr > __end;  
}
```

# The Worst Code I've Written?

```
__attribute__((__section__="__TEXT,__lcxx_override")))
void* operator new(size_t __size) noexcept {
    // ...
}
```

```
template <class _Ret, class... _Args>
bool __is_function_overridden(_Ret (*__fptr)(_Args...)) noexcept {
    extern char __segstart __asm("section$start$__TEXT$__lcxx_override");
    extern char __segend __asm("section$end$__TEXT$__lcxx_override");

    uintptr_t __start = (uintptr_t)(&__segstart);
    uintptr_t __end   = (uintptr_t)(&__segend);
    uintptr_t __ptr   = (uintptr_t)(__fptr);

    return __ptr < __start || __ptr > __end;
}
```

# The Worst Code I've Written?

```
__attribute__((__section__="__TEXT,__lcxx_override")))
void* operator new(size_t __size) noexcept {
    // ...
}
```

```
template <class _Ret, class... _Args>
bool __is_function_overridden(_Ret (*__fptr)(_Args...)) noexcept {
    extern char __segstart __asm("section$start$__TEXT$__lcxx_override");
    extern char __segend __asm("section$end$__TEXT$__lcxx_override");

    uintptr_t __start = (uintptr_t)(&__segstart);
    uintptr_t __end   = (uintptr_t)(&__segend);
    uintptr_t __ptr   = (uintptr_t)(__fptr);

    return __ptr < __start || __ptr > __end;
}
```

# The Worst Code I've Written?

```
__attribute__((__section__="__TEXT,__lcxx_override")))
void* operator new(size_t __size) noexcept {
    // ...
}
```

```
template <class _Ret, class... _Args>
bool __is_function_overridden(_Ret (*__fptr)(_Args...)) noexcept {
    extern char __segstart __asm("section$start$__TEXT$__lcxx_override");
    extern char __segend __asm("section$end$__TEXT$__lcxx_override");

    uintptr_t __start = (uintptr_t)(&__segstart);
    uintptr_t __end   = (uintptr_t)(&__segend);
    uintptr_t __ptr   = (uintptr_t)(__fptr);

    return __ptr < __start || __ptr > __end;
}
```

# Deployment Experience

- TMO adopted in the Darwin Kernel (XNU)
- Typed operator new adopted in Darwin user space system libraries
- Extremely effective
- Essentially no adoption cost

# Deployment Experience

- Not a silver bullet (not all allocations are funnelled through new)
- Effectiveness relies on QOI of the system allocator, which is a performance tradeoff

# Standardization Path

# Upcoming Paper

## P2719: Type-aware allocation and deallocation functions

Before

```
// user writes:  
new (args...) T(...)  
  
// compiler checks (in order):  
T::operator new(sizeof(T), args...)  
::operator new(sizeof(T), args...)
```

After

# Upcoming Paper

## P2719: Type-aware allocation and deallocation functions

Before

```
// user writes:  
new (args...) T(...)  
  
// compiler checks (in order):  
T::operator new(sizeof(T), args...)  
::operator new(sizeof(T), args...)
```

After

```
// user writes:  
new (args...) T(...)  
  
// compiler checks (in order):  
T::operator new(sizeof(T), args...)  
operator new(sizeof(T), type_identity<T>{}, args...)  
::operator new(sizeof(T), args...)
```

# Users Could Now Write

```
struct Druid : Character { };
struct Paladin : Character { };
struct Sorceress : Character { };

template <std::derived_from<Character> T>
void* operator new(std::size_t size, std::type_identity<T>) {
    // ... some special allocation scheme for these types ...
}
```

# A Conforming Extension Under the As-If Rule

```
template <class _Tp>
__attribute__((__overload_priority_(-1)))
void* operator new(std::size_t __size, std::type_identity<_Tp>) {
    std::__type_descriptor_t __descriptor = __builtin_type_descriptor(_Tp);
    // ... typed operator new implementation ...
}
```

# Conclusions

- Standard Library hardening tackles (mostly) spatial memory safety
  - Requires adoption to be effective
  - Great for bug finding and production “hardening”
  - Contracts provide a great framework for it
  - Go try it out!

# Conclusions

- TMO makes temporal memory issues harder to exploit
  - Adoption is almost 100% unintrusive
  - Does not fix any actual bugs, but makes them difficult to exploit
  - We propose a standardization path with other benefits

# Conclusions

- There's **a huge amount** of existing C++ code
  - A lot of it is unsafe by everyone's standard
  - We need to do something about that
  - Ease of adoption is a necessity

# Conclusions

- **Better** safety and security is achievable in C++
- We must look for simple and high impact changes, not perfection
- We encourage more WG21 work on immediate solutions
- Pragmatically consider the greater good, not only C++'s interests

# Thank You!