



How do Time Travel Debuggers Work?

Greg Law

**Most programmers
spend most of their
time debugging.**



*Everyone knows that debugging
is twice as hard as writing a
program in the first place.*

*So if you're as clever as you can
be when you write it, how will
you ever debug it?*

Brian Kernighan



THE #1 PROGRAMMER EXCUSE
FOR LEGITIMATELY SLACKING OFF:

"MY `printf`'s COMPILING."

HEY! GET BACK
TO WORK!

COMPILING!


OH. CARRY ON.

VS

STEP BY STEP DEBUGGING



How did that
happen?






unc



Expectations


Reality



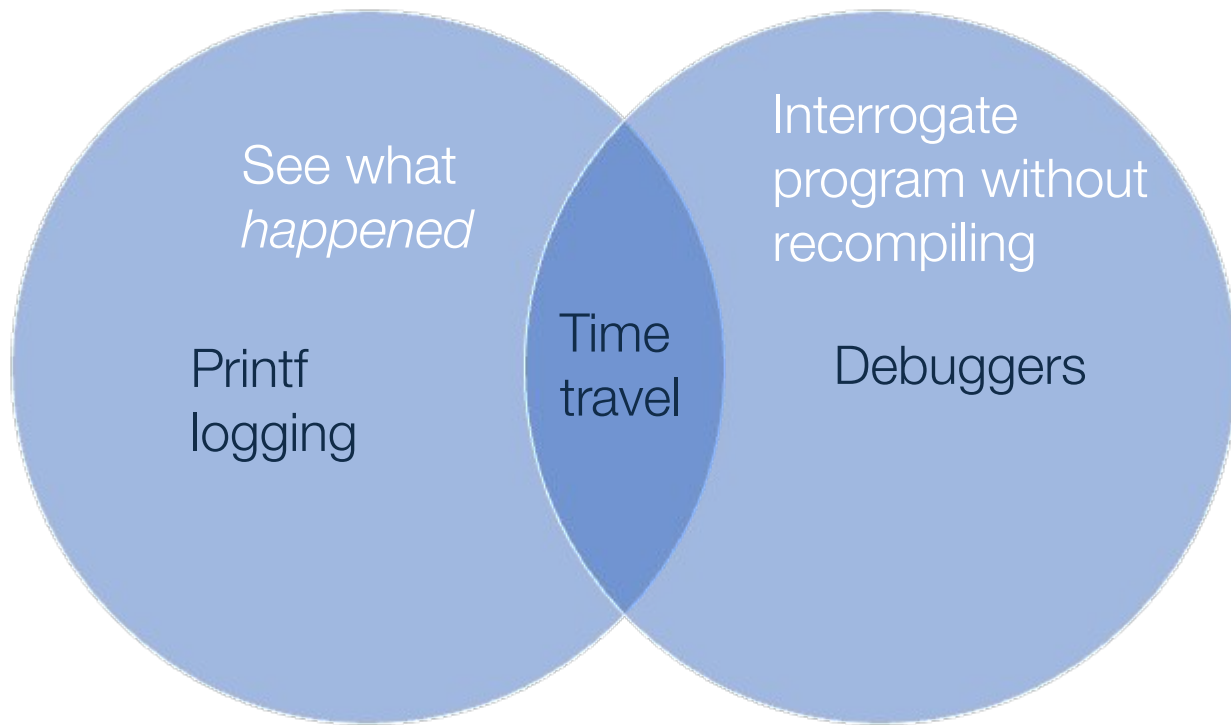




What makes bugs hard to fix?



What happened... in detail





C++ projects and products

Linux

- Undo - UDB & LiveRecorder
- rr (rr-project.org)
- GDB (ish)

Windows

- TTD

Embedded

- Lauterbach “TRACE32”
- Green Hills TimeMachine

undo[™]



Non C++

- JavaScript / React replay.io
- .Net RevDebug
Visual Studio
- Java Undo
- Rust, Go Undo / rr
- Python Undo*, PyPy*

Omniscient Debugger 29.Dec.06 - com.lambda.Debugger.Demo

File Run Trace Filter Previous Event 532 [1273] Demo.java:198

Threads

- <main_7>
- <Sorter_0>
- <Sorter_1>
- <Sorter_2>
- <Sorter_3>**
- <Sorter_4> --
- <Sorter_5> --
- <Sorter_6> --
- <Waiter_8> --

Stack

```
<DemoRunnable_3>.run()
<Demo_0>.sort(0, 5)
<Demo_0>.sort(3, 5)
<Demo_0>.average(3, 5)
```

Locals

* start	3
* end	5
* sum	0
* i	3

this

```
<Demo_0>
quick <Demo_1>
c 'X' (88)
b '=' (61)
array int[20]_0
```

Method Traces

```
***<DemoRunnable_3>.run() -> void
  <Demo_0>.sort(0, 5) -> void
    <Demo_0>.average(0, 5) -> 240
      DemoRunnable.new(<Demo_0>, 0, 2) -> <DemoRunnable>
      Thread.new(<DemoRunnable_6>, "Sorter") -> <Sorter>
      <Sorter_6>.start() -> void
      <Demo_0>.sort(3, 5) -> void
        <Demo_0>.average(3, 5) -> 483
          <Demo_0>.sort(3, 4) -> void
            <Demo_0>.sort(5, 5) -> void
              sort -> void
            <Sorter_6>.join() -> void
          sort -> void
        run -> void
      run -> void
```

Code

```
return;
}

public int average(int start, int end) {
    int sum = 0;
    for (int i = start; i < end; i++) {
        sum += array[i];
    }
}
```

TTY Output

```
-----ODB Demo Program-----
A badMethod threw: java.lang.NullPointerException.
Starting QuickSort: 20
-- Done sorting --
-- 0 1 --
-- 1 0 --
-- 2 237 --
-- 3 243 --
```

Objects

```
<Demo_0>
quick <Demo_1>
c 'X' (88)
b '=' (61)
array int[20]_0
* 19 1968
* 18 1962
17 1725
16 1719
* 15 1476
* 14 1470
13 1221
12 1233
11 1227
* 10 984
9 978
8 735
* 7 729
* 6 492
* 5 243
* 4 480
* 3 486
* 2 237
1 0
0 1
```

From last: 234 stamps, 0.017secs local = value

Design Decision 1.

How to remember

What was the previous state?

Two options:

1. Save it.
2. Recompute it.

$$a = a + 1$$

$$\rightarrow a = b$$

Deterministic re-execution

$$f(x) \rightarrow x'$$



Deterministic re-execution

- Snapshots.
- Event log.

Non-determinism

- What is unpredictable?
 - ~~System calls.~~
 - Thread switches.
 - Asynchronous events (signals).
 - Shared memory accesses.
 - Some machine instructions.

Non-deterministic code

```
0x00005555555515e  mov -0x4(%rbp),%eax
0x000055555555161  mov %eax,%esi
0x000055555555163  lea 0xe9a(%rip),%rdi
0x00005555555516a  mov $0x0,%eax
0x00005555555516f  sub %rdx,%rsp
0x000055555555171  mov %rcx,%r10
0x000055555555174  mov $0x11,%eax
0x000055555555179  syscall
```

Here's some machine code in a program that we want to be able to **record** and **replay**.

deterministic

nondeterministic

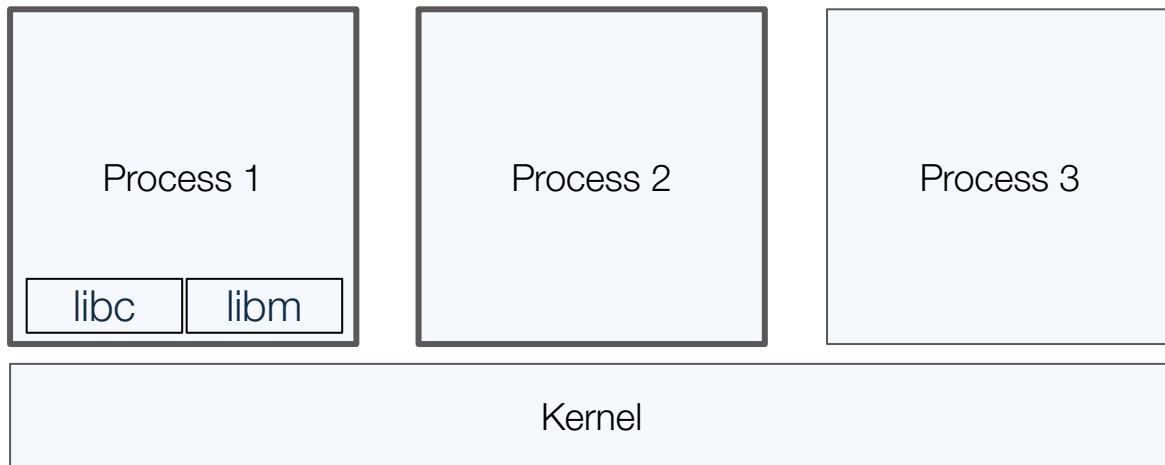
Design Decision 1

Re-execution is the winner

Design decision 2: what exactly to record?

- A process?
- A thread or function?
- A sandbox environment (e.g. JVM or v8?)
- Virtual machine?

Recording at process/OS ABI boundary



Design Decision 2 - winner!

Re-execution at the process boundary is the winner

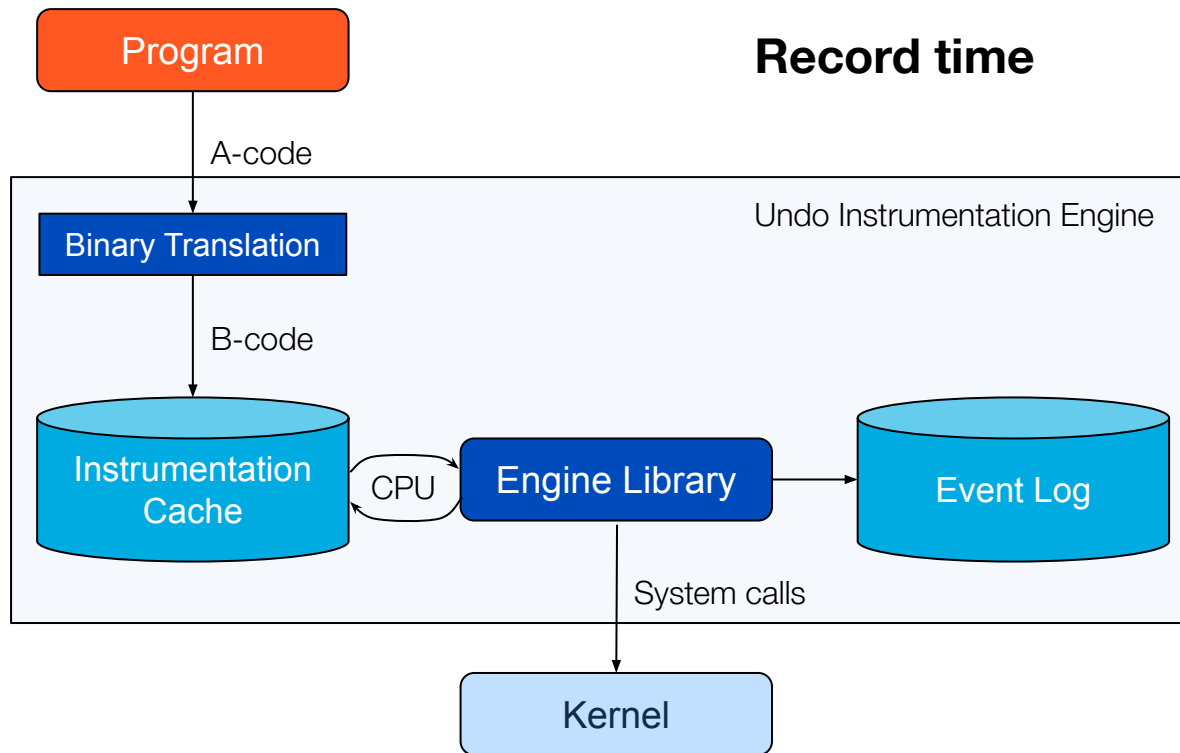
Design Decision 3: How to track time?



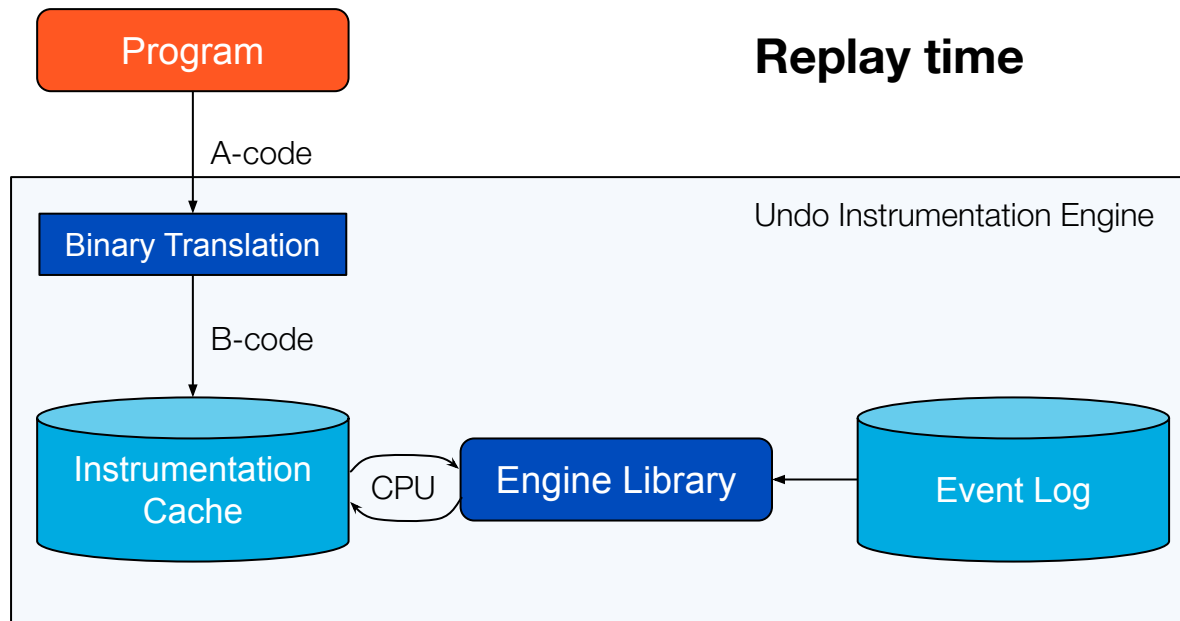
Design Decision 3: How to track time?

To JIT or not to JIT?

Translation and events



Translation and events



Perf counters

- Can be very fast.
- Simpler (less code).

JIT

- Works anywhere
- With anything*.
- Solid

Design Decision 4: Rely on memory determinism?


Memory Determinism

- (Much) smaller event logs
- Faster single-threaded

Not

- Unknown syscalls
- Parallel record

Thread serialization



Non-determinism

- What is unpredictable?
 - ~~System calls.~~
 - Thread switches.
 - Asynchronous events (signals).
 - ~~Shared memory accesses.~~
 - Some machine instructions.

Design decisions

- At what boundary to capture
- Binary rewriting instrumentation
- All/some/no memory recording
- Separate record/replay phases
- Parallel thread recording

Undo	rr	WinDbg	replay.io	ODB
proc	proc	proc	proc	JVM
yes	no	yes	no	no
some	none	all*	none	all
yes/no	yes	yes	yes	yes
no	no	yes	no	yes

Debuginfo

```
(gdb) print foo  
<value optimized out>  
(gdb)
```