


C++ now

# Implementing Ranges and Views

Roi Barkan



2024

# Hi, I'm Roi

- Roi Barkan (he/him) - רועי ברקן
- I live in Tel Aviv 🇮🇱
- C++ developer since 2000
- SVP Development & Technologies @ Istra Research
  - Finance, Low Latency, in Israel
  - [careers@istraresearch.com](mailto:careers@istraresearch.com)
- Always happy to learn and explore
  - Please - ask questions, make comments



[Slides](#)

# Outline

- Ranges and Views - Brief Intro
  - What are they
  - What's cool about them
  - Views we currently have
- Implementation Details - Several Perspectives
  - Object  $\Rightarrow$  Algorithm  $\Rightarrow$  Data
  - Concepts and Selection/Constraints
  - Lazy  $\Rightarrow$  Eager
- Case Study



# Ranges and Composition



# Ranges is a Breakthrough Library

- One of C++20 big-four features
- Rests on decades of existing libraries and experience
  - C++98 iterator-based algorithms
  - Fundamentals of functional / vectoric languages (APL, BQN, R, Julia, NumPy) [Conor Hoekstra](#)
  - Libraries of similar languages (D, Rust, Java) [Barry Revzin](#), [Alexandrescu BoostCon 2009](#).
- Main Innovation - Composability
  - Many algorithms take ranges as input and return ranges as output
    - Opposed to in-place or output-iterator nature of C++98 algorithms
  - Range Adaptors - algorithms encapsulated as 'lazy ranges' (views)
    - Algorithms as composable objects - 'expression templates'
  - Projections - unary transformations of the ranges we inspect.

# Terminology

- Range - Abstraction for a sequence of elements
  - begin-iterator and end-sentinel
- Range Algorithm - Function operating on ranges
  - Evolved from C++98 iterator based algorithms
  - Input: one or more ranges; potentially more arguments
  - Output: anything. If range: either in-place or via “output-iterator” or a subrange
- View - Ranges that are “cheap” to pass/hold
  - constant-time move, if-copyable-then-const-time (semantic nature → `enable_view<Rng>`)
- Range-Adaptor - range-to-range manipulations
  - Most adaptors are views and reside in `std::ranges::views`
  - View adaptors in the STL are ‘lazy’.
  - Adaptors are meant for chaining. The cheapness of views eases chain creation

# Composability of Ranges

- Chaining algorithms due to range arguments and results

```
ranges::reverse(ranges::search(str, "abc"sv)); godbolt
```

- Views as composable lazy ranges

```
str | views::split(' ') | views::take(2); godbolt
```

- Views have a value/algorithm duality

```
auto square_evens =  
    views::filter([](auto x) { return int(x) % 2 == 0; }) |  
    views::transform([](auto x) { return x * x; }); godbolt
```

- Simple combinations can enrich our vocabulary:

```
auto histogram =  
    views::chunk_by(std::equals{ }) |  
    views::transform([](const auto& rng) {  
        return make_pair(begin(rng), distance(rng));  
    });
```

# The Views in the Standard (C++20, C++23\*)

- Factories: `empty`, `single`, `iota`, `istream`, `repeat*`
- Rank preserving: `all`, `filter`, `transform`, `take{ _while }`, `drop{ _while }`, `subrange`, `counted`, `common`, `reverse`, `stride*`, `zip_transform*`, `adjacent_transform*`, `as_{const, rvalue}*`
- Rank decreasing - tuples: `elements`, `keys`, `values`
- Rank decreasing - ranges: `join{ _with* }`
- Rank increasing - tuples: `zip*`, `enumerate*`, `cartesian_product*`, `adjacent*`
- Rank increasing - ranges: `{lazy_}split`, `slide*`, `chunk{ _by }*`
- Committee plan for C++26 is in [P2760](#)





# Adaptor Chain Fundamentals



# Creating Composition Chains

- Adaptors support nesting as well as pipeline/infix composition
  - `views::take(views::split(str, ' '), 2)`  
equivalent to  
`str | views::split(' ') | views::take(2)` [godbolt](#)
- *RangeAdaptorClosure*: chains without a starting range
  - Objects that exist to be chained to some range
  - Semantically they are generic algorithms, not ranges
  - `std::ranges::range_adaptor_closure` is a CRTP helper for creating adaptors that have this nesting  $\Rightarrow$  pipeline duality.

# Simplest Range Adaptor

```
struct First : range_adaptor_closure<First> {  
    constexpr auto operator()(forward_range auto&& rng) const {  
        return subrange(begin(rng), empty(rng) ? begin(rng) : next(begin(rng)));  
    }  
};  
  
constexpr First first;  
  
int main() {  
    string s = "aa bb cc";  
    auto x = s | split(' ');  
    println("{} ", x | first);  
    return 0;  
}
```

Reference may dangle



```
Program returned: 0  
[['a', 'a']]
```

[godbolt](#)

# Dealing with Dangling

- Chains involve creation (and destruction) of temporary objects
- Solution - aggregate the chain into “expression templates”:
  - `typeid("x"s | split(' ') | take(3)) ≈ take_view<split_view<string>>`
- Adaptors themselves are typically small and cheap to pass as the chain grows
- Ranges can be expensive to pass → hence we use Views.

# Simplest Range Adaptor + View

```
template <view Inner> requires forward_range<Inner>
class FirstItemView : public view_interface<FirstItemView<Inner>> {
    [[no_unique_address]] Inner inner;
public:
    constexpr FirstItemView(Inner inner_) : inner(std::move(inner_)) {}
    constexpr auto begin() { return std::ranges::begin(inner); }
    constexpr auto end() { return empty(inner) ? begin() : next(begin()); }
    constexpr std::size_t size() { return empty(inner) ? 0 : 1; }
};

template <forward_range Range>
FirstItemView(Range&&) -> FirstItemView<views::all_t<Range>>;

struct First : range_adaptor_closure<First> {
    constexpr auto operator()(forward_range auto&& rng) const { return FirstItemView{rng}; }
}; godbolt
```

# Details About Views

- **view\_interface** - helper CRTP which opts-in to the **view** concept
- Constructor - pass inner view by-value, **std::move()** inside
- **begin()/end()** - must be implemented.
  - **const** correctness is tricky (see [Nico Josutis](#))
- **size()** - constant-time, opt-in a **sized\_range**.
- Deduction guide - use **views::all\_t** to allow non-view inputs
  - more about **all\_t** in the next slide
- Range adaptor closure - simply return the view.
  - Some adaptors can have optimizations here, e.g. **reverse** | **reverse**.

# Lifetime Management with **views::all**

- Chains of adaptors need to outlive their base range (otherwise UB).
- STL uses value categories (lvalue vs. rvalue) to try and avoid such cases
  - **ref\_view** - A view that points to another range (reference semantics), and cannot be constructed if the range is rvalue (about to go away)
  - **owning\_view** - A view that *takes ownership* of another range (moves it inside the view), and can be constructed only from rvalues. Move-only semantics (like **unique\_ptr**).
- **views::all(rng)** will return one of 3 different types of views:
  - If **rng** is a view - simply return it
  - else-if **rng** is a lvalue - return a **ref\_view** pointing to it (be careful of lifetimes)
  - else return an **owning\_view** that now owns the contents of the range.
- Range adaptor views in the STL use **views::all** to assist them.

# Examples - views, all

```
//temporaries create an owning view
static_assert( not view<decltype(string{" "})>> );
static_assert( view<decltype(string{" "})| all>> );
static_assert(is_same_v<decltype(string{" "})| all>,
              owning_view<string>.> );

//lvalues create a ref view
string s = "some string";
static_assert( not view<decltype(s)>> );
static_assert( view<decltype(s)| all>> );
static_assert(is_same_v<decltype(s)| all>,
              ref_view<string>.> ); godbolt
```



# Examples - views, all (2)

```
//views stay views
auto x = s | split(' ');
static_assert(    view<decltype(x)                )>);
static_assert(    view<decltype(x)                | all)>);
static_assert(is_same_v<decltype(x)                | all),
               decltype(x)                )>);

//Careful - all_t<array> can be expensive-to-move
static_assert( not view<decltype(array<int,1000>{}    )>);
static_assert(    view<decltype(array<int,1000>{} | all)>);
static_assert(is_same_v<decltype(array<int,1000>{} | all),
               owning_view<array<int,1000>>.    )>);
static_assert(sizeof(decltype(array<int,1000>{} | all)) >= 4000);godbolt
```

# Range Adaptor Iterators - Being Lazy

- Most views implement their own iterator (and/or sentinel) types, and achieve their functionality through the iterator member functions
  - `transform` - utilizing `operator*()`
  - `filter/stride/reverse` - utilizing `operator++()`
  - `take_while` - utilizing `operator!=(const sentinel&)`
  - `chunk/split` - utilizing `operator*()` and `operator++()`.
- The lazy approach has many benefits
  - Pay only for what you need
  - Better support for potentially infinite ranges
  - More data locality and less need for extra RAM
  - Compiler known expression-templates has potential for performance gains.

# See Barry About the Iterators

## Implementing filter in C++

```
template <input_range V, indirect_unary_predicate<iterator_t<V>> F>
class filter_view::Iterator {
    iterator_t<V> base_ = iterator_t<V>();
    filter_view* parent_ = nullptr;

public:
    using iterator_concept = /* ... */;
    using iterator_category = /* ... */;
    using reference = range_reference_t<V>;
    using value_type = range_value_t<V>;
    using difference_type = range_difference_t<V>;


    Iterator() = default;
    Iterator(filter_view&, iterator_t<V>);

    auto operator*() const -> reference {
        return *base_;
    }

    auto operator++() -> Iterator& {
        base_ = find_if(++base_, ranges::end(parent_>base_), parent_>fun_);
        return *this;
    }

    auto operator++(int) -> Iterator {
        auto tmp = *this;
        ++this;
        return tmp;
    }

    auto operator==(Iterator const& rhs) const -> bool {
        return base_ == rhs.base_;
    }
};
```



Barry Revzin



CODE RECKONS

Science to the CORE

CPPP21

Take(5)

# Range Categories and Refinements

- Ranges are categorized by their power of iteration, similar to the C++98 iterator category model
  - input → forward → bidirectional → random-access → contiguous
  - Similarly to C++98 category is associated via opt-in of `iterator_category` tags.
- On top of the power of iteration, ranges have additional *orthogonal* refinements:
  - borrowed - iterators can outlive the range. opt-in **`enable_borrowed_range`**
  - sized - number of elements in amortized constant time. opt-out **`disable_sized_range`**
  - common - **`begin()`** and **`end()`** return the same type
  - constant - range into read-only values.
- Range Adaptors must correctly publish their effect on their input.

# Refinement Motivation - Algorithm Selection

- Sometimes the same goal can be achieved in several ways
  - `ranges::ssize` - returns a signed integer equal to the size of a range
  - `ranges::distance` - returns the distance between the beginning and end of a range
  - `ssize` only works for *sized* ranges (constant-time calculation)  
`distance` allows linear calculation if necessary. [Ben Deane recommends it.](#)
- The library uses concepts to constrain which ranges are applicable for which algorithm/view, and to know the best method of reaching the intended goal
- Before C++20 other mechanisms were used to achieve this goal - and with concepts we have a way to be more precise and more flexible where needed.

# Digression - How Lazy are We

- Recall histogram. How many passes does it perform over the data

```
auto histogram =  
    views::chunk_by(std::equals{ }) |  
    views::transform([](const auto& rng) {  
        return make_pair(begin(rng), distance(rng));  
    });
```

- Intuitively a single pass is enough.
- Depends on if `range_reference_t<chunk_by_view<...>>` is sized
  - i.e. depends on if `subrange<...>` is sized.
  - Could potentially be controlled via `subrange_kind` but not possible in existing adaptors
- Alternative implementation can `enumerate` and properly chunk/transform the pairs.

# Range/Iterator **const** Correctness

- Remember that iterators have indirect semantics.
- Still, ranges were meant to differentiate between **iterator** and **const\_iterator** for 'deep' constness.
- Views are thus allowed to differentiate and have 2 different iterator types.
- C++23 now has **std::basic\_const\_iterator** which can be used as a drop in iterator adaptor.
- Views are notoriously tricky (bad) when it comes to const-correctness
  - Due to caching behavior
  - Due to **owning\_view** vs. **ref\_view** being so interchangeable
  - See [Nico Josutis](#).

# Iterator Customization Points

- Apart of the basic operators (`*`, `!=`, `++`, `-`, `+=`, ...), iterators are allowed implement two more functions, which the ranges library must use for their purpose:
- `iter_move(iterator)` - instead if `std::move(*iterator)`
- `iter_swap(it1, it2)` - instead if `std::ranges::swap(*it1, *it2)`
- Main motivation: proxy-iterators (e.g. `zip_view`)
  - More on that from [Jacob Rice](#).
- Typically implemented as “hidden friends” and invoked via `std::ranges::iter_{move, swap}` - which are CPOs



# CPO - Customization Point Objects

- Customization points - ways in which a library (ranges) allows its users (specific range-adaptor implementers) to dictate how it behaves in certain cases.
- Before C++20 the STL had “clunky” customization point mechanisms
  - Template specialization (e.g. `std::hash`) [\[unord.hash\]](#)
  - Overload resolution and ADL (e.g. `std::swap`) [\[swappable.requirements\]](#).
- CPOs are actually objects (global variables) with template `operator()` function which knows to perform the correct search for customized implementations (typically via `if constexpr` or `requires` clauses)
  - More on that from [Gašper Ažman](#).



# Case Study



# Views for Sorted Ranges ([More Ranges Please](#))

- Suggestion - views for `merge`, `set_union`, `set_intersection`, `set_{symmetric}_difference`
  - Most algorithms can benefit from multi-input implementations
  - Heap (`priority_queue`) is needed for efficient `set_union`, `merge`, ....
- STL contains several algorithms for sorted ranges: `{inplace}_merge`, `includes`, `set_{union,intersection,{symmetric}_difference}`
  - Also search algorithms: `{upper,lower}_bound`, `equal_range`, `(unique)`.
- All the operations are lazy in nature
- Ranges-v3 [has views](#) for `set_{union,intersection,{symmetric}_difference}` with 2 input ranges
- D-lang has [merge](#) and [multiWayMerge](#).

# Implementation Approach

- Every STL algorithm with an output-iterator result can be conceptually converted to a lazy range-adaptor view.
- Basic approach - the unified iterator holds all sub-iterators, an indication of the 'current' one and a pointer to the range.
  - Key idea is that every call to `operator++()` should iteratively increment the lowest sub-iterator until a condition (based on the specific algorithm) is satisfied.
- Various details and opportunities exist for the different algorithms

# Set Operation Details

- **begin()** in constant-time
  - Trivial for union, merge. Caching needed for intersection, difference.
- Iterator category
  - input iteration seems enough (single pass)
  - forward/bidirectional iteration can be preserved - bidirectional needs a second heap.
  - random-access on either input can be utilized, mostly for intersection and difference (e.g. `lower_bound`)
  - random-access cannot be preserved.
- **common\_range** can be preserved.
- **sized\_range** can be preserved for merge.

# Set Operations on Multiple Inputs

- Variadic (compile time) input-count should be simple
  - Potentially use `array<variant<iterator_t<Views>...>, sizeof...(Views)>` with heap operations like `make_heap`, `pop_heap`, `push_heap`.
- Dynamic Range-of-Ranges is more tricky due to potential RAM needs.

Potential approaches:

- Take a random-access container as extra argument.
- Take a (PMR) allocator as extra argument.
- Expect the input range (of ranges) to be random-access and use it (like D-lang [multiWayMerge](#))

```
auto carsByPrice =
    carsByMakerThenPrice | chunk_by([](const Car& a, const Car& b) {
                                    return a.maker == b.maker;
    }) | to<vector> |
    merge([](const Car& a, const Car& b) { return a.price < b.price; });
```

# Alternative Approach - `std::generator`

- C++23's first library addition utilizing coroutines.
- A **generator** exposes a coroutine with `co_yield` calls as a **view**.
- Main advantage - simplicity:
  - All the intermediate state can be stored in variables
  - Procedural style instead of callback style
  - I don't think one **generator** can be implemented for all output-iterator range algorithms - "the coloring problem".
- Main disadvantages:
  - Exposes an `input_range`, not more
  - Performance is compiler/optimizer dependent.

# Summary

- The C++ ranges library is an exemplar of composability
- Ranges were developed to be enhanced and extended
- Implementing ranges code requires know-how
  - Not rocket science
- Now it's our turn
  
- Thank you !!
  - Questions and comments are welcome

[Slides](#)

