

C++ now

<Random> Harder Better Faster Stronger

Adrien Devresse

2024

<Random>
Harder, Better, Faster
Stronger



About me

Adrien Devresse

- Around 15 years of C++
- **Background** in High Performance Computing, Distributed systems, Embedded systems.
- **Worked at**
 - CERN
 - Human Brain Project
 - Startup
 - Woven By Toyota
 - (Onboarding) Square Kilometer Array

Disclaimer

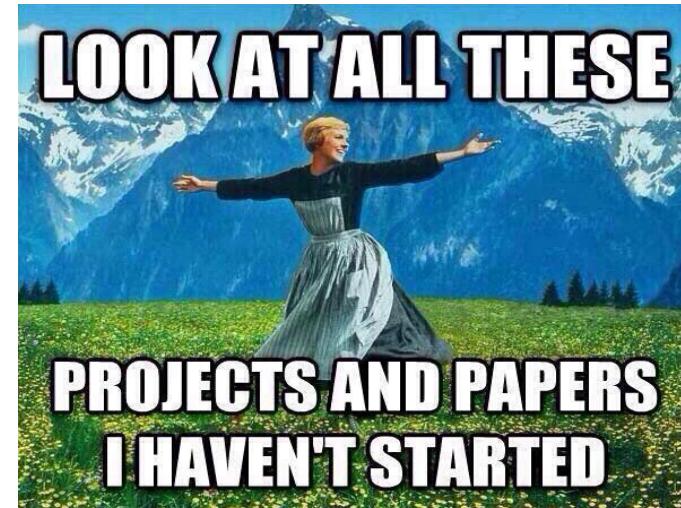
This talk is about:

- **Practical** usages of pseudo random number generators. Not the theory behind them
- Centered on **C++11 <Random>**
- Aims to non-cryptographic applications **only**
- **Focus** on Scientific Computing

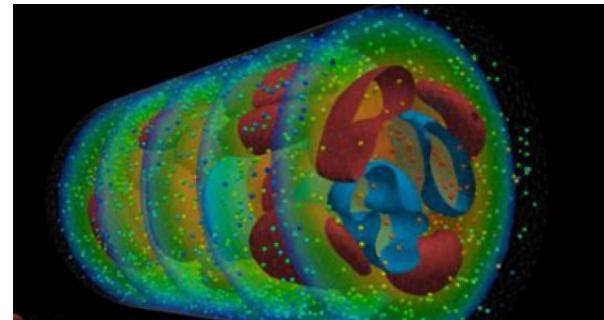
The little story of Bob



Bob
*Average Phd
Student*



Procrastinate



**Write a scientific simulator to finish
his Phd**

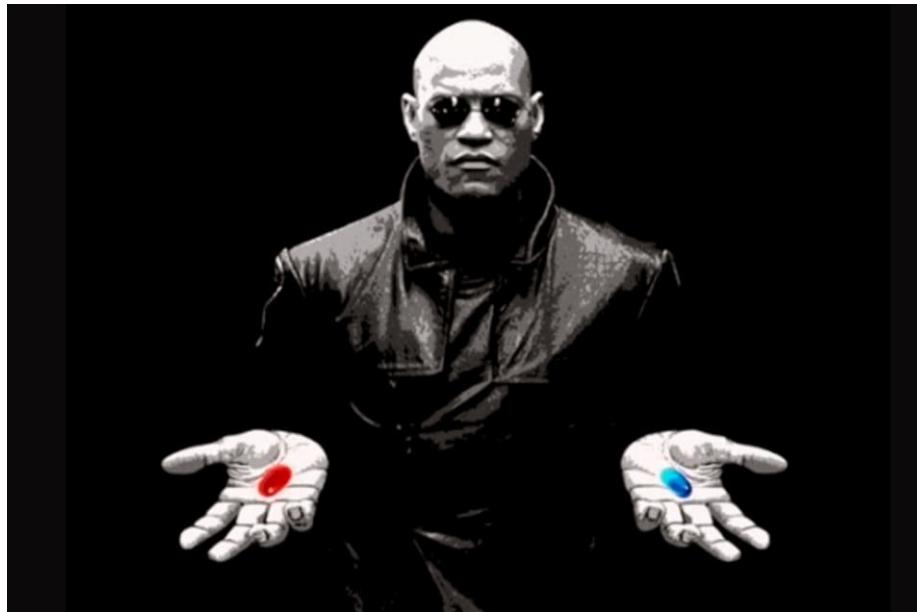
Few things about Scientific Computing

- Most of Scientific Simulators are written in 
- They need to be **fast**
- They rely heavily on **Random number generation**
- They need the random streams to have **excellent randomness** properties
- They use extensively **parallel processing**

Easy

Choose your weapon ~~weapon~~ **PRNG**

C
Λ<~~r~~>d>



C++

<random>

Harder

C

```
// Seed  
std::srand()  
  
// Generate  
std::rand()
```

C++

minstd_rand0 (C++11)	std::linear_congruential_engine<std::uint_fast32_t, 16807, 0, 2147483647> Discovered in 1969 by Lewis, Goodman and Miller, adopted as "Minimal standard" in 1988 by Park and Miller
minstd_rand (C++11)	std::linear_congruential_engine<std::uint_fast32_t, 48271, 0, 2147483647> Newer "Minimum standard", recommended by Park, Miller, and Stockmeyer in 1993
mt19937 (C++11)	std::mersenne_twister_engine<std::uint_fast32_t, 32, 624, 397, 31, 0x9908b0df, 11, 0xffffffff, 7, 0x9d2c5680, 15, 0xefc60000, 18, 1812433253> 32-bit Mersenne Twister by Matsumoto and Nishimura, 1998
mt19937_64 (C++11)	std::mersenne_twister_engine<std::uint_fast64_t, 64, 312, 156, 31, 0xb5026f5aa9619e9, 29, 0x5555555555555555, 17, 0x71d67ffffeda6000, 37, 0xffff7eee00000000, 43, 6364136223846793005> 64-bit Mersenne Twister by Matsumoto and Nishimura, 2000
ranlux24_base (C++11)	std::subtract_with_carry_engine<std::uint_fast32_t, 24, 10, 24>
ranlux48_base (C++11)	std::subtract_with_carry_engine<std::uint_fast64_t, 48, 5, 12>
ranlux24 (C++11)	std::discard_block_engine<std::ranlux24_base, 223, 23> 24-bit RANLUX generator by Martin Lüscher and Fred James, 1994
ranlux48 (C++11)	std::discard_block_engine<std::ranlux48_base, 389, 11> 48-bit RANLUX generator by Martin Lüscher and Fred James, 1994
knuth_b (C++11)	std::shuffle_order_engine<std::minstd_rand0, 256>
default_random_engine (C++11)	implementation-defined

Non-deterministic random numbers

std::random_device is a non-deterministic uniform random bit generator, although implementations are allowed to implement std::random_device using a pseudo-random number engine if there is no support for non-deterministic random number generation.

random_device (C++11) non-deterministic random number generator using hardware entropy source (class)

Random number distributions

A random number distribution post-processes the output of a URNG in such a way that resulting output is distributed according to a defined statistical probability density function.

Random number distributions satisfy *RandomNumberDistribution*.

Defined in header <random>

Uniform distributions

uniform_int_distribution (C++11)	produces integer values evenly distributed across a range (class template)
uniform_real_distribution (C++11)	produces real values evenly distributed across a range (class template)

Bernoulli distributions

bernoulli_distribution (C++11)	produces <code>bool</code> values on a Bernoulli distribution (class)
binomial_distribution (C++11)	produces integer values on a binomial distribution (class template)
negative_binomial_distribution (C++11)	produces integer values on a negative binomial distribution (class template)
geometric_distribution (C++11)	produces integer values on a geometric distribution (class template)

Poisson distributions

poisson_distribution (C++11)	produces integer values on a Poisson distribution (class template)
exponential_distribution (C++11)	produces real values on an exponential distribution (class template)
gamma_distribution (C++11)	produces real values on a gamma distribution (class template)
weibull_distribution (C++11)	produces real values on a Weibull distribution (class template)
extreme_value_distribution (C++11)	produces real values on an extreme value distribution (class template)

Normal distributions

normal_distribution (C++11)	produces real values on a standard normal (Gaussian) distribution (class template)
lognormal_distribution (C++11)	produces real values on a lognormal distribution (class template)
chi_squared_distribution (C++11)	produces real values on a chi-squared distribution (class template)

Not that bad



PRNG



Statistical Distributions



minstd_rand0 (C++11)	std::linear_congruential_engine<std::uint_fast32_t, 16807, 0, 2147483647> Discovered in 1969 by Lewis, Goodman and Miller, adopted as "Minimal standard" in 1988 by Park and Miller
minstd_rand (C++11)	std::linear_congruential_engine<std::uint_fast32_t, 48271, 0, 2147483647> Newer "Minimum standard", recommended by Park, Miller, and Stockmeyer in 1993
mt19937 (C++11)	std::mersenne_twister_engine<std::uint_fast32_t, 32, 624, 397, 31, 0x9908b0df, 11, 0xffffffff, 7, 0x9d2c5680, 15, 0xefc60000, 18, 1812433253> 32-bit Mersenne Twister by Matsumoto and Nishimura, 1998
mt19937_64 (C++11)	std::mersenne_twister_engine<std::uint_fast64_t, 64, 312, 156, 31, 0xb5026f5aa9619e9, 29, 0x5555555555555555, 17, 0x71d67ffffeda6000, 37, 0xffff7eee00000000, 43, 6364136223846793085> 64-bit Mersenne Twister by Matsumoto and Nishimura, 2000
ranlux24_base (C++11)	std::subtract_with_carry_engine<std::uint_fast32_t, 24, 10, 24>
ranlux48_base (C++11)	std::subtract_with_carry_engine<std::uint_fast64_t, 48, 5, 12>
ranlux24 (C++11)	std::discard_block_engine<std::ranlux24_base, 223, 23> 24-bit RANLUX generator by Martin Lüscher and Fred James, 1994
ranlux48 (C++11)	std::discard_block_engine<std::ranlux48_base, 389, 11> 48-bit RANLUX generator by Martin Lüscher and Fred James, 1994
knuth_b (C++11)	std::shuffle_order_engine<std::minstd_rand0, 256>
default_random_engine (C++11)	implementation-defined

Non-deterministic random numbers

std::random_device is a non-deterministic uniform random bit generator, although implementations are allowed to implement std::random_device using a pseudo-random number engine if there is no support for non-deterministic random number generation.

random_device (C++11) non-deterministic random number generator using hardware entropy source

Random number distributions

A random number distribution post-processes the output of a URNG in such a way that resulting output is distributed according to a defined statistical probability density function.

Random number distributions satisfy *RandomNumberDistribution*.

Defined in header <random>

Uniform distributions

uniform_int_distribution (C++11)	produces integer values evenly distributed across a range (class template)
uniform_real_distribution (C++11)	produces real values evenly distributed across a range (class template)

Bernoulli distributions

bernoulli_distribution (C++11)	produces bool values on a Bernoulli distribution (class)
binomial_distribution (C++11)	produces integer values on a binomial distribution (class template)
negative_binomial_distribution (C++11)	produces integer values on a negative binomial distribution (class template)
geometric_distribution (C++11)	produces integer values on a geometric distribution (class template)

Poisson distributions

poisson_distribution (C++11)	produces integer values on a Poisson distribution (class template)
exponential_distribution (C++11)	produces real values on an exponential distribution (class template)
gamma_distribution (C++11)	produces real values on a gamma distribution (class template)
weibull_distribution (C++11)	produces real values on a Weibull distribution (class template)
extreme_value_distribution (C++11)	produces real values on an extreme value distribution (class template)

Normal distributions

normal_distribution (C++11)	produces real values on a standard normal (Gaussian) distribution (class template)
lognormal_distribution (C++11)	produces real values on a lognormal distribution (class template)
chi_squared_distribution (C++11)	produces real values on a chi-squared distribution (class template)

PRNGs more in details

	Quality	Speed	Periodicity	Age
Ranlux			10^{166}	1994
Mersenne Twister			10^{19937}	1998-2003
Minstd			Short	1988-1993
Knutb			Short	1988



std::default_random_engine -> Minstd

More on Mersenne Twister

- Developed by [Makoto Matsumoto](#) and [Takuji Nishimura](#)
 - *Mersenne Twister: A 623-Dimensionally Equidistributed Uniform Pseudo-Random Number Generator*
- Gold standard
- Implemented in many languages



Lets see how it looks Sequential

```
std::uint64_t seed = 42;  
  
std::mt19937_64 rng(seed);  
  
for(std::size_t i= 0; i< dataset.size(); i++){  
    std::uniform_int_distribution dist;  
    auto number = dist(rng);  
    my_kernel(number, dataset[i] );  
}
```

Lets see how it looks **Parallel**

Version 0

```
std::uint64_t seed = 42;  
  
std::mt19937_64 rng(seed);  
std::mutex mut;  
  
#pragma omp parallel for  
for(std::size_t i= 0; i< dataset.size(); i++){  
    std::uniform_int_distribution dist;  
    auto number = [](){  
        std::unique_lock guard(mut);  
        return dist(rng);  
    };  
    my_kernel(number, dataset[i] );  
}
```



Lets see how it looks **Parallel**

Version 1

```
std::uint64_t seed = 42;  
thread_local std::mt19937_64 rng;  
thread_local bool rng_initialized = false;
```

#pragma omp parallel for

```
for(std::size_t i= 0; i< dataset.size(); i++){  
    std::uniform_int_distribution dist;  
    If(!rng_initialized ){  
        rng.seed(seed * total_thread_numbers + thread_id);  
        rng_initialized = true;  
    }  
    auto number = dist(rng);  
    my_kernel(number, dataset[i] );  
}
```

WRONG

Lets see how it looks **Parallel**

Version 2

```
std::uint64_t seed = 42;  
  
thread_local std::mt19937_64 rng(seed);  
thread_local bool rng_initialized = false;  
  
#pragma omp parallel for  
for(std::size_t i= 0; i< dataset.size(); i++){  
    std::uniform_int_distribution dist;  
    If(!rng_initialized ){  
        rng.discard( (1 << 32 ) * thread_id);  
        rng_initialized = true;  
    }  
    auto number = dist(rng);  
    my_kernel(number, dataset[i] );  
}
```

Better

A Solution ?

$$f(k, d) = R$$
$$P(R) \equiv N$$

f: Function

k: Key

d: Data

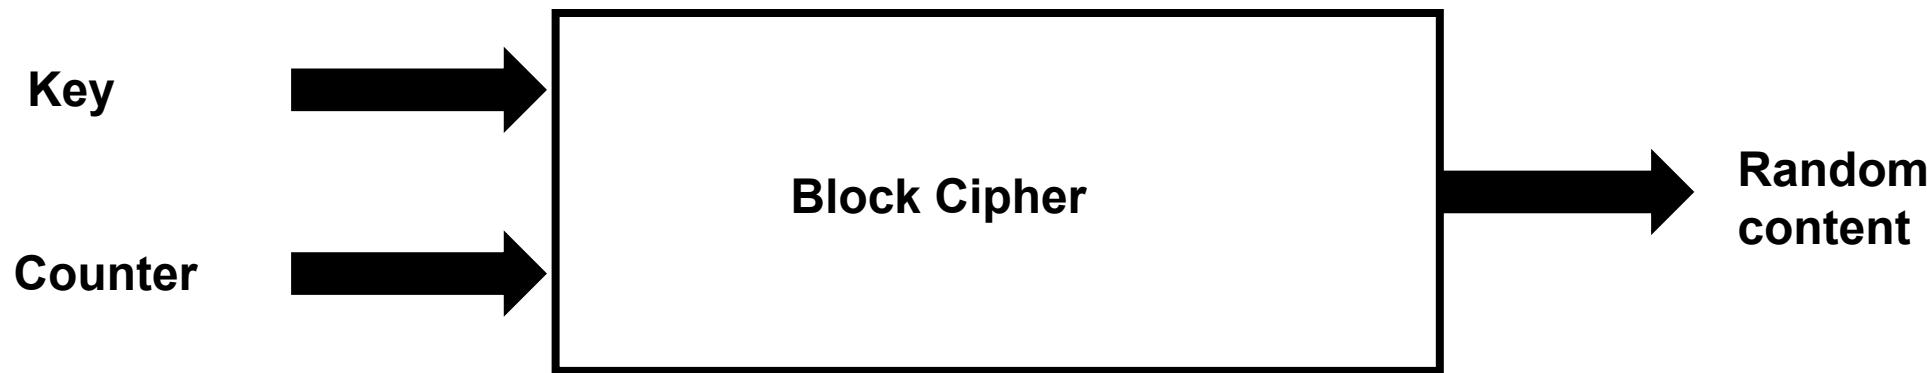
R: Result

N: Random Noise



Cryptographic Block Cipher

Counter Based Pseudo Random Number Generator (CBRNG)



N_0 : Key=seed, C=0

N_1 : Key=seed, C=1

N_2 : Key=seed, C=2

- No problem to derivate the seed in parallel algorithms
 - **K = tuple(seed, computing_bloc_id, ...)**
- C is known and predictable
 - **Our random generator can be stateless**

Back to how it looks

Version CBRNG

```
std::uint64_t seed = 42;  
constexpr std::uint64_t block_size = 1024;
```

```
#pragma omp parallel for  
for(std::size_t i= 0; i < dataset.size(); i++){  
    std::uniform_int_distribution dist;  
    alea::counter_rng<threefry> gen;  
    gen.set_key({ seed, i / block_size });  
    get.set_counter( i % block_size );  
    auto number = dist(rng);  
    my_kernel(number, dataset[i] );  
}
```

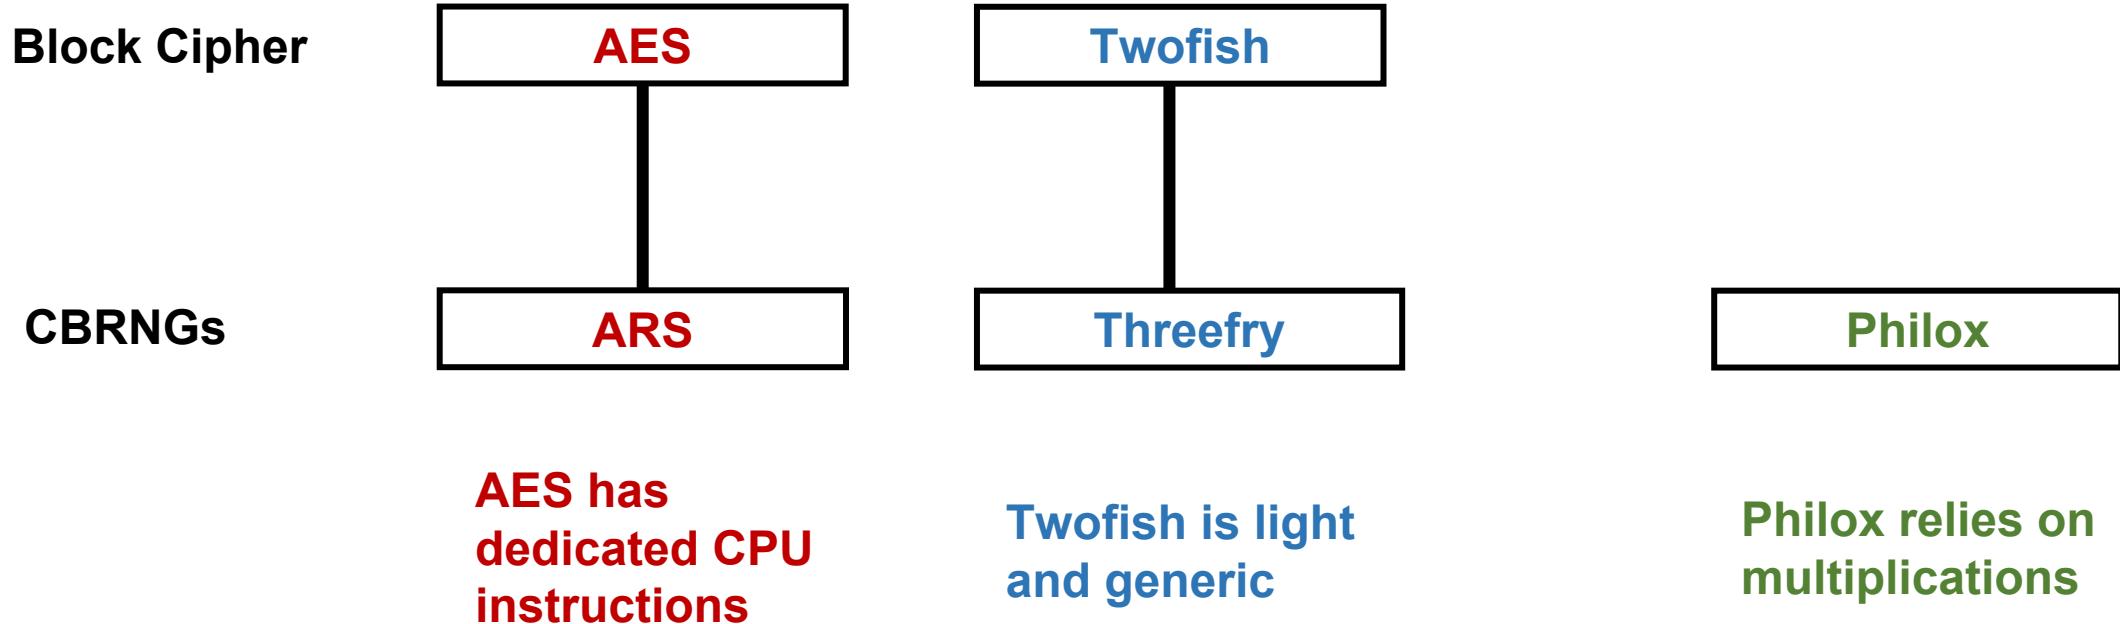


Faster

Random123 to the rescue

- Publication *Parallel Random Numbers: As Easy as 1, 2, 3*,
Salmon, Moraes, Dror & Shaw, **Supercomputing, 2011**
- Define a entire family of **CBRNG**
- Based (Mostly) on battle-tested **Block cipher algorithms**
- **Reduced number of rounds** for a speed tradeoff

Random123: 3 CBRNGs



Performances

PRNG	CPU (Intel) GB/s	GPU (Nvidia) GB/s	GPU (AMD) GB/s
Mersenne Twister	6.1	18.3	N/A
ARS (With Hardware instructions)	11.1	N/A	N/A
Threefry	6.4	52.8	46.4
Philox	3.4	145.3	79.1

Stronger

Random Quality

- Quality of Random generator is evaluated based on their statistical properties
- The reference test suite for that named **TestU01**
 - Publication: *TestU01: A C library for empirical testing of random number generators*
- The test suit is divided into two components:
 - Crush (96 tests)
 - Big Crush (106 tests)

TestU01 conformance

PRNGs	Crush compliance	Big Crush compliance
Mersenne Twister	Full	Partial
ARS	Full	Full
Threefry	Full	Full
Philox	Full	Full

Conclusions

- `<random>` is an amazing piece of code within the STL **but**
 - It aged (C++11)
 - The defaults shall (probably) be revisited
 - It can be **hard** to use right
- CBRNGs provide **better** algorithms for Many-cores applications
- CBRNGs can often be **faster** than most `<random>` algorithms
- **Random123** algorithms provide **stronger** randomness quality than most `<random>` algorithms

Opinion: I believe we shall update the PRNGs algorithms within the STL. **A good candidate for a Beman project ?**



Twitter: <https://twitter.com/adev>

Mastodon: @adev@hostux.social

Discord: adev9145