

C++ now

Reflection is Good for (Code) Health

Saksham Sharma

2024

YOURS TRULY

- Quantitative research technology at Tower Research Capital
 - High frequency trading firm based out of NYC
- Develop low latency trading systems (C++)
 - Nanoseconds and microseconds
- Develop high throughput research systems (C++ and Python)
 - $O(\text{terabytes})$ data
- Program analysis and functional programming in a past life
- Love performance, software abstractions, and clean APIs

THE BORING PART

This talk's contents are mine and mine alone

- Not my employer's :)

OVERVIEW

- What is reflection
- Reflection in other languages (Go, Python, Java)
- Reflection in C++ as per P2996
 - Syntax and examples
- Reflection libraries!
 - Python bindings
 - ABI hashing (`boost::abi_hash?`)
 - A duck-typed `std::any` (`boost::virtual_any?`)
- Alternatives ways to achieve “reflection”

REFLECTION?

In code.



REFLECTION?

Ability to write code in a language such that:

- Access information about other “code” in a programmatic form and operate on it

```
class MyClass {  
    int a;  
    int b;  
};  
  
for (auto member_info : gimme_class_members<MyClass>()) {  
    std::cout << "member - " << member_info.name() << std::endl;  
}
```

REFLECTION?

HOW IS THIS DIFFERENT FROM METAPROGRAMMING / TEMPLATES?

REFLECTION?

- Not much :)
- Operate on `constexpr` values instead of types
 - Syntactic difference, not semantic
- Compiler can provide richer “information” about code
 - Some of it is already there
 - A neat and consistent “bag” of features to expand in the future

RUNTIME REFLECTION

Ability to access information about other “code” at runtime

- Could be done in interpreted or compiled languages
 - But generally seen in interpreted languages

```
class MyClass:
    def __init__(self):
        self.x = 1
        self.y = 2

for member_name, member_value in MyClass().__dict__.items():
    print(f"{member_name}: {member_value}")
```

COMPILE TIME REFLECTION / STATIC REFLECTION

Ability to access information about other “code”:

- At compile time

```
class MyClass {  
    int a;  
    int b;  
};  
  
for (auto member_info : gimme_class_members<MyClass>()) {  
    if (is_ptr(member_info)) { ... }  
}
```

EXAMPLES FROM OTHER LANGUAGES

Venturing into some
alien worlds



MY FAVORITE EXAMPLE: PYTHON

At runtime your code can:

- Modify other code to do entirely different things when a method is called.
- Say you run some code and it can change what a method on your object actually does.
- Change what it means to access a field on an object.
- Add new methods or attributes to any object.

FUN WITH REFLECTION - PYTHON

```
def modify_cls(cls):
    if not hasattr(cls, "copy"):
        return cls
    orig_copy = cls.copy
    def _wrapped_copy(obj):
        print("Calling wrapped copy")
        attrs = obj.__dict__.keys()
        print("Attributes: " +
              " ".join(attrs))
        result = orig_copy(obj)
        return result
    cls.copy = _wrapped_copy
```

```
class MyClass:
    def __init__(self, x):
        self.x = x
    def copy(self):
        return MyClass(self.x)

>>> modify_cls(MyClass)
>>> MyClass(2).copy()
Calling wrapped copy
Attributes: x
<__main__.MyClass object at 0x7f1a9e6>
```

GOLANG

- Golang is a compiled but duck-typed language
 - Well, structurally typed, but close enough
- Relies heavily on interfaces
- Runtime reflection similar to python.
 - No special compile time constructs
 - Uses the `reflect` package's methods to get “reflection values”.

DUCK TYPING?



DUCK.GO

```
type Vehicle interface {
    Start() string
}

type Car struct {
    Make string
}

func (this Car) Start() { // Car::Start(Car* this)
    return "Brrr"
}

type Truck struct {
    WheelCount int
}

func (this Truck) Start() {
    return "Brhhhhhhhhhhh"
}
```

```
func (v Vehicle) SelfDrive() {
    fmt.Println(v.Start())
}

func main() {
    vehicles := []Vehicle{
        Car{"Toyota"},
        Truck{4},
    }
    for _, v := range vehicles {
        SelfDrive(v)
    }
}
```


FUN WITH REFLECTION - GOLANG

```
type T struct {  
    A int  
    B string  
}  
  
t := T{23, "skidoo"}  
s := reflect.ValueOf(&t).Elem()  
typeOfT := s.Type()  
for i := 0; i < s.NumField(); i++ {  
    f := s.Field(i)  
    fmt.Printf("%d: %s %s = %v\n", i, typeOfT.Field(i).Name, f.Type(), f.Interface())  
}  
  
// 0: A int = 23  
// 1: B string = skidoo
```

*<https://go.dev/blog/laws-of-reflection>

JAVA

- Surprisingly, quite similar in feel to Python and Go.
- Reflection is “runtime”, in the sense that the object type saves type information accessible at program runtime.
- `java.lang.reflect`

FUN WITH REFLECTION - JAVA

```
Class cls = Class.forName("method1"); // Surprising, lookup types with string!
Method methlist[] = cls.getDeclaredMethods();
for (int i = 0; i < methlist.length; i++) {
    Method m = methlist[i];
    System.out.println("name = " + m.getName());
    System.out.println("decl class = " + m.getDeclaringClass());
}
Class pvec[] = m.getParameterTypes();
for (int j = 0; j < pvec.length; j++)
    System.out.println("param #" + j + " " + pvec[j]);
```

SCARY?



ZERO COST ABSTRACTIONS

- Prior examples had “reflection objects” at runtime.
- Despite the temptation, we must do compile time
 - No one likes the cost of RTTI
- We can lean on constexpr algorithms and metaprogramming!
 - Compile time programming is turing complete after all :)

A NEW HOPE!

P2996



WE HAVE A SYNTAX AND AN IMPLEMENTATION!

- P2996 is a promising paper gaining traction.
 - Wyatt Childers, Peter Dimov, Dan Katz, Barry Revzin, Andrew Sutton, Faisal Vali, Daveed Vandevoorde
- Two working implementations already!
 - Edison Design Group (EDG) compiler on Godbolt
 - Clang fork by Bloomberg on Godbolt and GitHub
 - Godbolt link for both: <https://godbolt.org/z/cGK4Eo6K1>

WE HAVE A SYNTAX AND AN IMPLEMENTATION!

- Good consensus on syntax and some semantics.
- Value based reflection is an interesting choice
 - We'll talk about it soon :)
- Prior work: Reflection TS
 - David Sankel - Type based reflection

WE HAVE A REFLECTION OPERATOR

unary operator ^

“Lifts” into reflection land

i.e. produces a reflection value

- function
- variable and friends
- non-static data member
- template
- constant expression
- namespace
- ...

```
auto magic_reflection_method(std::meta::info obj) {  
    .. do something with “meta” info ..  
}
```

```
auto result = magic_reflection_method(^MyType);
```

WE HAVE A `STD::META::INFO`

A formless type that describes all meta-information about a type / member / method / etc.

```
auto magic_reflection_method(std::meta::info obj) {  
    .. do something with “meta” info ..  
}  
  
auto result = magic_reflection_method(^MyType);
```

WE HAVE A **SPLICE** OPERATOR

`[: r :] =>`

Takes a `std::meta::info`

Constant expression

Splices it back into your

regular code

```
struct MyStruct {  
    static int a;  
    static int b;  
};  
  
template <typename T>  
void get_member(std::meta::info elem) {  
    std::cout << T::[:elem:] << std::endl;  
}  
  
std::cout  
    << get_member<MyStruct>(^MyStruct::a)  
    << std::endl;
```

WE HAVE A `STD::META::DEFINE_CLASS`

We can create classes in thin air! (almost)

Need to declare it first, need to name all the members and their types.

```
struct S;  
static_assert(is_type(define_class(^S, {  
    data_member_spec(^int, {  
        .name="i", .align=64  
    })),  
    data_member_spec(^int, {  
        .name="j", .align=64  
    })),  
    ));
```

WE HAVE A `STD::META::DEFINE_CLASS`

Equivalent of....


```
struct S;  
struct S {  
    alignas(64) int i;  
    alignas(64) int j;  
};
```

WE HAVE A LOT OF CONSTEVAL METHODS TO GO WITH IT

```
namespace std::meta {  
    // [meta.reflection.names], reflection names and locations  
    consteval string_view name_of(info r);  
    consteval string_view qualified_name_of(info r);  
    ...  
    consteval bool is_function(info r);  
    consteval bool is_variable(info r);  
    consteval bool is_type(info r);  
    // [meta.reflection.member.queries], reflection member queries  
    template<class... Fs>  
    consteval vector<info> members_of(info type, Fs... filters);  
    ...  
}
```

ELEVATOR PITCH

We really do need this
in C++!



ENUM TO STRING

```
enum class MyEnum { VALUE_1, VALUE_2, MAX_VALUES };

template <typename EnumT> constexpr std::string enum_to_string(EnumT enum_value) {
    template for (constexpr auto e : std::meta::enumerators_of(^EnumT)) {
        if (enum_value == [:e:]) return std::string(std::meta::name_of(e));
    }
    return "<unnamed>";
}

template <typename EnumT> constexpr EnumT string_to_enum(std::string enum_str) {
    // some blackmagic
}
```


EASY CMDLINE PARSING

```
struct MyOpts {  
    std::string file_name = "input.txt"; // Option "--file_name <string>"  
    int count = 1;                       // Option "--count <int>"  
};  
  
int main(int argc, char* argv[]) {  
    MyOpts opts = parse_options<MyOpts>(  
        std::vector<std::string_view>(argv + 1, argv + argc)  
    );  
}
```

EASY CMDLINE PARSING

```
template <typename Opts> auto parse_options(ArgT args) -> Opts {
    Opts opts;
    template for (constexpr auto dm : nonstatic_data_members_of (^Opts)) {
        auto it = std::ranges::find_if(args, ... match string ...);
        auto iss = std::ispanstream(it[1]);
        if (iss >> opts.[:dm:]; !iss) {
            std::print(stderr, "Failed to parse option {}\n", *it);
            std::exit(EXIT_FAILURE);
        }
    }
    return opts;
}
```

BETTER CMDLINE PARSING

```
struct Args : Clap {
    Option<std::string, {.use_short=true, .use_long=true}> name;
    Option<int, {.use_short=true, .use_long=true}> count = 1;
};

int main(int argc, char** argv) {
    auto opts = Args{}.parse(argc, argv);
    for (int i = 0; i < opts.count; ++i) { // opts.count has type int
        std::print("Hello {}!", opts.name); // opts.name has type std::string
    }
}
```

OPERATING ON A CLASS' LAYOUT

```
constexpr auto get_layout() {  
    constexpr auto members = nonstatic_data_members_of (^S);  
    std::array<my_descriptor, members.size()> layout;  
    for (int i = 0; i < members.size(); ++i) {  
        layout[i] = {.offset = offset_of(members[i]), .size = size_of(members[i])};  
    }  
    return layout;  
}
```

ARRAY OF STRUCTS TO STRUCT OF ARRAYS

```
struct point {  
    float x;  
    float y;  
};  
  
using points = struct_of_arrays<point, 30>;  
  
// equivalent to:  
// struct points {  
//     std::array<float, 30> x;  
//     std::array<float, 30> y;  
// };
```

WHY THIS
'FORMLESS'
STD::META::INFO?

TYPE VS VALUE BASED REFLECTION

- Type based reflection is more intuitive
 - `std::meta::method` / `std::meta::variable` / `std::meta::class`
 - Can imagine writing template specializations etc
 - Maybe ranges and algorithms make life better with value based?

TYPE VS VALUE BASED REFLECTION

- Type based reflection is more intuitive
 - `std::meta::method` / `std::meta::variable` / `std::meta::class`
 - Can imagine writing template specializations etc
 - Maybe ranges and algorithms make life better with value based?
- Slower to compile
 - Lot more types to handle for the compiler

TYPE VS VALUE BASED REFLECTION - COMPILE TIMES

```
constexpr int f() {  
    int i = 0;  
    for (int k = 0; k < 10000; ++k)  
        i += k;  
    return i / 10000;  
}  
  
template <int N> struct S {  
    static constexpr int sm = S<N - 1>::sm + f();  
};  
  
template <> struct S<0> {  
    static constexpr int sm = 0;  
};  
  
constexpr int r = S<200>::sm;
```

```
struct Int { int v; };  
  
constexpr int f() {  
    Int i = {0};  
    for (Int k = {0}; k.v < 10000; ++k.v)  
        i.v += k.v;  
    return i.v / 10000;  
}  
  
template <int N> struct S {  
    static constexpr int sm = S<N - 1>::sm + f();  
};  
  
template <> struct S<0> {  
    static constexpr int sm = 0;  
};  
  
constexpr int r = S<200>::sm;
```

TYPE VS VALUE BASED REFLECTION

- Type based reflection is more intuitive
 - `std::meta::method` / `std::meta::variable` / `std::meta::class`
 - Can imagine writing template specializations etc
 - Maybe ranges and algorithms make life better with value based?
- Slower to compile
 - Lot more types to handle for the compiler
- Makes language **rigid** by preventing future breakages
 - Major issue, downstream code can start depending on which type a specific reflection would be.
 - Sometimes might not be clear which type to classify a reflection into

TYPE VS VALUE BASED REFLECTION

- We want users to write “duck typed” reflection code
 - They check if that reflection value has the feature they want from it
 - If so, you can use it in your desired context.
- Flexible because every usage would *ideally* be gated by a feature check
 - Easier to add feature checks compared to renaming meta classes
- Feature checks have better cross-interaction than inheritance

WHAT DO WE KNOW SO FAR?

Language	Runtime or compile-time	Type or value based
Python	Runtime	?
Golang	?	?
Java	?	?
C++26 🙌	Compile-time	Value (std::meta::info)

LET'S TABULARIZE THIS

Language	Runtime or compile-time	Type or value based (loosely)
Python	Runtime	Value (with a sprinkle of Type)
Golang	Runtime	Value (with a sprinkle of Type)
Java	Runtime	Type
C++26 🙌	Compile Time	Value (std::meta::info)

HOW DO WE USE THIS NEW SUPERPOWER ?

No one likes chaos.



LIBRARIES!

- Makes it easier to write general-purpose / boilerplate-reducing libraries
- Libraries would require fewer redundant inputs from the user
 - A CLI parsing library won't need you to enumerate all objects in your struct each time.
 - Maybe it could take a spec to parse and return a struct to you storing that spec?

LIBRARIES!

PYTHON BINDINGS, ABI HASHING, A BETTER STD::ANY

PYTHON BINDINGS

A FAN FAVORITE, AND ONE OF THE MOST COMMON WISHLIST ITEMS

- We have a C++ class that we'd like to expose to Python
- Background:
 - You can run python and C++ code in the same process.
 - CPython is a C library and application.
 - Can run in the same process that is running your C++ code
 - We have to define how to “expose” C++ objects in python

```
struct Item {  
    int id;  
    PyObject* getPyValue() const;  
};  
  
struct Row {  
    const auto& items() const { return items_; }  
    auto nanotime() const { return nanotime_; }  
};  
  
class RowReader : public BinaryListener {  
public:  
    RowReader(const std::string& filename)  
        : reader_(filename);  
    std::string getIdName(int id) const;  
    const auto& getRows();  
};
```

A FAN FAVORITE, AND ONE OF THE MOST COMMON WISHLIST ITEMS

- We don't want to write *this*
- Most of the code is just repetition of what we already know
- Some of the things are non-obvious
 - Whether a function's return type should be wrapped as a reference or a value copy
 - How to name a type in Python

```
BOOST_PYTHON_MODULE(binary_reader_bpy) {  
    class_<Item>("Item", no_init)  
        .def_readonly("id", &Item::id)  
        .def("value", &Item::getPyValue);  
  
    class_<Row, std::shared_ptr<Row>>("Row", no_init)  
        .def("nanotime", &Row::nanotime)  
        .def("items", &Row::items,  
            return_internal_reference<>());  
  
    class_<RowReader, boost::noncopyable>(  
        "RowReader", init<std::string>(arg("filename"))  
    )  
        .def("getRows", &RowReader::getRows)  
        .def("getIdName", &RowReader::getIdName);  
}
```

REFLECTION MAKES IT EASY PEASY

```
template <typename T> object make_python_type() {
    std::string cls_name{meta::name_of (^T)};
    auto type_obj = class_<T>(cls_name.c_str(), init<int, int>());
    [:expand(meta::members_of (^T)):] >> [&]<auto e> {
        if constexpr(!meta::is_public(e)) { return; }
        std::string name{meta::name_of(e)};
        if constexpr(meta::is_nonstatic_data_member(e))
            type_obj.def_readwrite(name.c_str(), &[:e:]);
        if constexpr(meta::is_function(e) && !meta::is_constructor(e) && !meta::is_destructor(e)) {
            if constexpr(!std::is_reference_v<typename return_type<decltype(&[:e:])>::type>)
                type_obj.def(name.c_str(), &[:e:]);
        }
    };
    return type_obj;
}
```

KIND OF...

DECLARE THE TYPE OBJECT

```
template <typename T> object make_python_type() {
    std::string cls_name{meta::name_of (^T)};
    auto type_obj = class_<T>(cls_name.c_str(), init<int, int>());
    [:expand(meta::members_of (^T)):] >> [&]<auto e> {
        if constexpr(!meta::is_public(e)) { return; }
        std::string name{meta::name_of(e)};
        if constexpr(meta::is_nonstatic_data_member(e))
            type_obj.def_readwrite(name.c_str(), &[:e:]);
        if constexpr(meta::is_function(e) && !meta::is_constructor(e) && !meta::is_destructor(e)) {
            if constexpr(!std::is_reference_v<typename return_type<decltype(&[:e:])>::type>)
                type_obj.def(name.c_str(), &[:e:]);
        }
    };
    return type_obj;
}
```

LOOP OVER THE MEMBERS AND METHODS

```
template <typename T> object make_python_type() {
    std::string cls_name{meta::name_of (^T)};
    auto type_obj = class_<T>(cls_name.c_str(), init<int, int>());
    template for (auto e : meta::members_of (^T)) {
        if constexpr(!meta::is_public(e)) { return; }
        std::string name{meta::name_of(e)};
        if constexpr(meta::is_nonstatic_data_member(e))
            type_obj.def_readwrite(name.c_str(), &[:e:]);
        if constexpr(meta::is_function(e) && !meta::is_constructor(e) && !meta::is_destructor(e)) {
            if constexpr(!std::is_reference_v<typename return_type<decltype(&[:e:])>::type>)
                type_obj.def(name.c_str(), &[:e:]);
        }
    };
    return type_obj;
}
```

LOOP OVER THE MEMBERS AND METHODS

```
template <typename T> object make_python_type() {
    std::string cls_name{meta::name_of (^T)};
    auto type_obj = class_<T>(cls_name.c_str(), init<int, int>());
    [:expand(meta::members_of (^T)):] >> [&]<auto e> {
        if constexpr(!meta::is_public(e)) { return; }
        std::string name{meta::name_of(e)};
        if constexpr(meta::is_nonstatic_data_member(e))
            type_obj.def_readwrite(name.c_str(), &[:e:]);
        if constexpr(meta::is_function(e) && !meta::is_constructor(e) && !meta::is_destructor(e)) {
            if constexpr(!std::is_reference_v<typename return_type<decltype(&[:e:])>::type>)
                type_obj.def(name.c_str(), &[:e:]);
        }
    };
    return type_obj;
}
```


PROBABLY DON'T WANT PUBLIC MEMBERS EXPOSED

```
template <typename T> object make_python_type() {
    std::string cls_name{meta::name_of (^T)};
    auto type_obj = class_<T>(cls_name.c_str(), init<int, int>());
    [:expand(meta::members_of (^T)):] >> [&]<auto e> {
        if constexpr(!meta::is_public(e)) { return; }
        std::string name{meta::name_of(e)};
        if constexpr(meta::is_nonstatic_data_member(e))
            type_obj.def_readwrite(name.c_str(), &[:e:]);
        if constexpr(meta::is_function(e) && !meta::is_constructor(e) && !meta::is_destructor(e)) {
            if constexpr(!std::is_reference_v<typename return_type<decltype(&[:e:])>::type>)
                type_obj.def(name.c_str(), &[:e:]);
        }
    };
    return type_obj;
}
```

INFORM BOOST::PYTHON HOW TO ACCESS THE MEMBERS??

```
template <typename T> object make_python_type() {
    std::string cls_name{meta::name_of (^T)};
    auto type_obj = class_<T>(cls_name.c_str(), init<int, int>());
    [:expand(meta::members_of (^T)):] >> [&]<auto e> {
        if constexpr(!meta::is_public(e)) { return; }
        std::string name{meta::name_of(e)};
        if constexpr(meta::is_nonstatic_data_member(e))
            type_obj.def_readwrite(name.c_str(), &MyStruct::a);
        if constexpr(meta::is_function(e) && !meta::is_constructor(e) && !meta::is_destructor(e)) {
            if constexpr(!std::is_reference_v<typename return_type<decltype(&[:e:])>::type>)
                type_obj.def(name.c_str(), &[:e:]);
        }
    };
    return type_obj;
}
```

INFORM BOOST::PYTHON HOW TO ACCESS THE MEMBERS

```
template <typename T> object make_python_type() {
    std::string cls_name{meta::name_of (^T)};
    auto type_obj = class_<T>(cls_name.c_str(), init<int, int>());
    [:expand(meta::members_of (^T)):] >> [&]<auto e> {
        if constexpr(!meta::is_public(e)) { return; }
        std::string name{meta::name_of(e)};
        if constexpr(meta::is_nonstatic_data_member(e))
            type_obj.def_readwrite(name.c_str(), &[:e:]);
        if constexpr(meta::is_function(e) && !meta::is_constructor(e) && !meta::is_destructor(e)) {
            if constexpr(!std::is_reference_v<typename return_type<decltype(&[:e:])>::type>)
                type_obj.def(name.c_str(), &[:e:]);
        }
    };
    return type_obj;
}
```

INFORM BOOST::PYTHON ABOUT METHODS

```
template <typename T> object make_python_type() {
    std::string cls_name{meta::name_of (^T)};
    auto type_obj = class_<T>(cls_name.c_str(), init<int, int>());
    [:expand(meta::members_of (^T)):] >> [&]<auto e> {
        if constexpr(!meta::is_public(e)) { return; }
        std::string name{meta::name_of(e)};
        if constexpr(meta::is_nonstatic_data_member(e))
            type_obj.def_readwrite(name.c_str(), &[:e:]);
        if constexpr(meta::is_function(e) && !meta::is_constructor(e) && !meta::is_destructor(e)) {
            if constexpr(!std::is_reference_v<typename return_type<decltype(&[:e:])>::type>)
                type_obj.def(name.c_str(), &[:e:]);
        }
    };
    return type_obj;
}
```

IT WORKS :)

```
template <typename T> object make_python_type() {
    std::string cls_name{meta::name_of (^T)};
    auto type_obj = class_<T>(cls_name.c_str(), init<int, int>());
    [:expand(meta::members_of (^T)):] >> [&]<auto e> {
        if constexpr(!meta::is_public(e)) { return; }
        std::string name{meta::name_of(e)};
        if constexpr(meta::is_nonstatic_data_member(e))
            type_obj.def_readwrite(name.c_str(), &[:e:]);
        if constexpr(meta::is_function(e) && !meta::is_constructor(e) && !meta::is_destructor(e)) {
            if constexpr(!std::is_reference_v<typename return_type<decltype(&[:e:])>::type>)
                type_obj.def(name.c_str(), &[:e:]);
        }
    };
    return type_obj;
}
```

PYTHON BINDINGS

- Can be done
- Still need to figure out how to customize things that can't be defaulted properly
 - Return types with reference
 - Picking overloads for functions
- How to name types (and methods in case of overloads)
- Docstrings

CUSTOMIZING THE DEFAULT BEHAVIOR

- Defining template specializations!
- Not ideal, but has benefits
 - Can annotate even if you don't have control on source code

```
constexpr auto customizations = {  
    {^Row::items,  
     return_value_policy::reference_internal},  
    ...  
};
```

USER DEFINED ATTRIBUTES

- Proposed in P1887, discussed in P2911
- Helpful in tagging information at the place of definition
- Requires control on the source code of the type

```
class Row {  
public:  
    [[return_policy("reference_internal")]]  
    const auto& items() { ... }  
};
```


USER DEFINED ATTRIBUTES

- One of my favorite features from Golang
- Annotating at the point of definition
 - Has its place, as opposed to annotating at the time of use

```
type User struct {  
    Name string `json:"name" required:"true"`  
}  
  
user := User{"John"}  
field, ok := reflect.TypeOf(user).Elem().FieldByName("Name")  
fmt.Println(field.Tag, field.Tag.Get("required"))  
// json:"name" required:"true" true
```

*<https://www.makeuseof.com/reflection-in-go/>

SUMMARY

- Default behavior can be done easily
- Customizations are trickier / not DRY
- User defined attributes might help!
 - In serialization / deserialization / CLI parsing too!

ABI HASHING

Why and how

SETTING THE STAGE - WHY IS A TYPE'S HASH USEFUL?


- Say you had a unique hash for a type's memory layout.
 - What would you do?
- Example message sent between two modules communicating.

SchemaType	Rest of your data
------------	-------------------

- We'd like to minimize the space we're using to describe the data's schema in a given message.
- Ideally good to keep the comparison fast as well.

SETTING THE STAGE - WHY IS A TYPE'S HASH USEFUL?

- Say you had a unique hash for a type's memory layout.
 - What would you do?



SETTING THE STAGE - WHY IS A TYPE'S HASH USEFUL?

- Say you had a unique hash for a type's memory layout.
 - What would you do?
- Why not compare the “type” in text form?
 - Situations where speed is important
 - Cannot compare full type layouts each time for speed
 - Handshakes can help
 - Situations where size is important
 - Messages written to disk or sent over the network

SETTING THE STAGE - WHY IS A TYPE'S HASH USEFUL?

- Say you had a unique hash for a type's memory layout.
 - What would you do?
- Send this over UDP for handshake-less connections
 - Connections with handshakes can use a full description of the layout
 - That's hard enough already without protobuf etc.
 - Sending over a hash is reasonably size-efficient

SETTING THE STAGE - WHY IS A TYPE'S HASH USEFUL?

- Say you had a unique hash for a type's memory layout.
 - What would you do?
- Use it to safely type-cast across module boundaries
 - Multiple python modules?
 - Different .so files talking to each other with non-trivial types at the ABI boundary.

SETTING THE STAGE - WHY IS A TYPE'S HASH USEFUL?

- Say you had a unique hash for a type's memory layout.
 - What would you do?

libcreator.so

```
struct MyCls {  
    int x;  
    virtual int foo() = 0;  
};  
shared_ptr<MyCls> create() {  
    return make_shared<Impl>(22);  
}  
  
class_<MyCls, shared_ptr<MyCls>>(  
    "MyClsPy", no_init  
);
```

```
obj = creator.create()  
# type(obj) == creator.MyClsPy  
  
usermodule.use(obj)  
# ERROR!  
# usePy expects  
#   usermodule.MyClsPy
```

libusermodule.so

```
struct MyCls {  
    int x;  
    virtual int foo() = 0;  
};  
int use(shared_ptr<MyCls> obj) {  
    return obj->foo() + obj->x;  
}  
  
class_<MyCls, shared_ptr<MyCls>>(  
    "MyClsPy", no_init  
);
```

ISSUES WITH A SIMPLE VERSION FIELD

- Doesn't handle quirks that humans may miss (eg: attributes)
- Easy to forget to change version
 - Especially when transitive dependencies are changed
- Can't do it generically in a library, user has to handle

```
struct BaseT {  
    int a;  
} __attribute__((packed));  
  
struct MyMessage : public BaseT {  
    int b;  
};
```

OKAY HOW DO WE HASH A TYPE'S LAYOUT?

REFLECTION!!

REFLECTION TO ITERATE A TYPE'S LAYOUT

- Is a decent test of the capabilities of reflection as proposed in P2996.
- Requires recursively computing the hash of types of member variables of a class. Avoid cycles!
- Requires a constexpr hashing function.
- Requires full visibility into the class' internals
 - Sounds scary actually, we can now write code that can know about private members of a class

REFLECTION TO ITERATE A TYPE'S LAYOUT

- Also, requires some level of configurability

```
struct ABIHashingConfig {  
    static constexpr int MINIMUM_SUPPORTED_VERSION = 0; // To allow future rollover  
    static constexpr int MAXIMUM_SUPPORTED_VERSION = 0; // To gracefully error out  
    uint8_t version : 4 = 0;  
    bool include_nsdmm_names : 1 = true;  
    bool include_indirections : 1 = false; // Only relevant in intra-process  
} __attribute__((packed));
```

ABI HASHING - THE TEST CASE

```
struct Order { int side = 1; size_t quantity = 0; };  
template <typename T> struct MyList {  
    T* start = nullptr;  
    T* end = nullptr;  
    bool valid : 1 = false;  
};  
template <typename T> struct MyList2 : public MyList<T> {};  
struct OrderBook {  
    MyList2<Order> buy_orders;  
    MyList2<Order> sell_orders;  
};
```

ABI HASHING - THE API

```
template <typename T, std::meta::abi::ABIHashingConfig config>  
constexpr size_t get_abi_hash();
```

```
// ...or...
```

```
constexpr size_t get_abi_hash(std::meta::info R,  
                              std::meta::abi::ABIHashingConfig config);
```

ABI HASHING - RECURSE OVER BASE CLASSES

```
constexpr size_t _get_abi_hash_impl(  
    std::meta::info R,  
    meta::abi::ABIHashingConfig config = meta::abi::ABIHashingConfig{}) {  
    size_t hash = 0;  
    for (auto e : bases_of(R)) {  
        hash = hash_combine(hash, get_abi_hash(meta::type_of(e), config));  
    };  
};
```


ABI HASHING - LOOPING OVER MEMBERS - IGNORING POINTERS

```
for (auto e : nonstatic_data_members_of(R)) {  
    auto elem_type = meta::type_of(e);  
    auto is_indirect_ref = (meta::test_type (^std::is_pointer_v, elem_type) ||  
                            meta::test_type (^std::is_reference_v, elem_type));  
    if (!config.include_indirections && is_indirect_ref) {  
        // Maybe even warn or throw, since no use-case.  
        continue;  
    }  
}
```

ABI HASHING - LOOPING OVER MEMBERS - MEMBERS' LAYOUT

```
for (auto e : nonstatic_data_members_of(R)) {  
    if (config.include_nsdms_names) {  
        auto name = std::string(meta::name_of(e));  
        hash = hash_combine(hash, HASH_STR(name.c_str()));  
    }  
    if (meta::is_bit_field(e)) {  
        hash = hash_combine(hash, std::meta::offset_of(e),  
                             std::meta::bit_offset_of(e), std::meta::bit_size_of(e));  
    } else {  
        hash = hash_combine(hash, std::meta::offset_of(e), std::meta::size_of(e));  
    }  
}
```

ABI HASHING - LOOPING OVER MEMBERS - TYPE OF THE MEMBER?

```
for (auto e : nonstatic_data_members_of(R)) {  
    if (meta::test_type (^std::is_class_v, elem_type)) {  
        hash = hash_combine(hash, get_abi_hash(meta::type_of(e), config));  
    } else {  
        if (meta::test_type (^std::is_pointer_v, elem_type)) {  
        } else if (meta::test_type (^std::is_reference_v, elem_type)) {  
        } else {  
            auto name = std::string(meta::name_of(meta::type_of(e)));  
            hash = hash_combine(hash, HASH_STR(name.c_str()));  
        }  
    }  
}
```

ABI HASHING - LOOPING OVER MEMBERS - TYPE OF THE MEMBER?

```
for (auto e : nonstatic_data_members_of(R)) {  
    if (meta::test_type (^std::is_class_v, elem_type)) {  
    } else {  
        if (meta::test_type (^std::is_pointer_v, elem_type)) {  
            if (config.include_indirections) {  
                hash = hash_combine(hash, HASH_STR("pointer"));  
                constexpr auto ctype =  
                    std::meta::substitute(^std::remove_pointer_t, {elem_type});  
                hash = hash_combine(hash, get_abi_hash(ctype, config));  
            }  
        }  
    }  
}
```

ABI HASHING - FINALLY, SIZE

```
// Despite all members being the size, attributes like `packed` may change
// the size of the struct. Not everyone would be concerned with padding at
// the end though. Can consider making this optional via a config param.
if (std::meta::size_of(R) > 0) {
    hash = hash_combine(hash, std::meta::size_of(R));
}
```

ABI HASHING - AVOIDING CYCLES

```
constexpr size_t get_abi_hash(std::meta::info R,  
                               meta::abi::ABIHashingConfig config = meta::abi::ABIHashingConfig{},  
                               std::vector<std::meta::info> active_types = {}) {  
    auto it = std::ranges::find_if(active_types, [R](const auto& elem) { return elem == R; });  
    if (it == active_types.end()) {  
        active_types.push_back(R);  
        return _get_abi_hash_impl(R, config, active_types);  
    } else {  
        // Cycle detected! Return the index since we must modify the hash still.  
        return (it - active_types.begin());  
    }  
}
```

AND JUST FOR FUN, THIS MIGHT HAVE LOOKED LIKE THIS WITH TYPE BASED CODE

ABI HASHING - AVOIDING CYCLES

```
template <typename T, ABIHashingConfig config, typename... ActiveTypes>
constexpr size_t get_abi_hash() {
    constexpr ssize_t type_index = mp11::mp_find<mp_list<ActiveTypes...>, T>();
    if constexpr(type_index != mp11::mp_size<mp_list<ActiveTypes...>>()) {
        return static_cast<size_t>(type_index);
    } else {
        return _get_abi_hash_impl<T, config, ActiveTypes...>();
    }
}
```


ABI HASHING - SO WHAT DO WE HAVE NOW

- A simple way to ensure compatibility of data layout across processes, without using protobuf etc
- A solution that solves 80% of the common set of requirements of users
 - Can a python binding library cast one type to the other safely?
 - Are these two networked binaries using the same data layout?
- Requires minimal to no work on the user's ends
- Inflexible and hard to modify / handle unique situations

ABI HASHING - WHERE DO WE GO FROM HERE?

- Full ABI textual representation!
 - Or, an ABI for describing the ABI
- You could generate a “schema” file for your structs
 - Could be some JSON based schema
 - Could be a pre-existing schema like Apache Avro
 - This way we get a full ecosystem (with cross language support) for free.

A PYTHONIC STD::ANY

Hot take incoming



BOOST::PYTHON HAD A DREAM

```
object f(object x, object y) {  
    if (y == "foo")  
        x.slice(3,7) = "bar";  
    else  
        x.attr("items") += y(3, x);  
    return x;  
}  
  
// Duck typing and a completely untyped type!
```

STD::ANY HAD A DREAM

```
// any type
std::any a = 1;
std::cout << a.type().name() << ": " << std::any_cast<int>(a) << '\n';
a = 3.14;
std::cout << a.type().name() << ": " << std::any_cast<double>(a) << '\n';
a = true;
std::cout << a.type().name() << ": " << std::any_cast<bool>(a) << '\n';

// No duck typing sadly, but we got the equivalent of std::variant<everything>
// Can't do much though :)
```



**BOOST::VIRTUAL_ANY
WITH DUCK TYPING**

STD::ANY

**BOOST::PYTHON
::OBJECT**

HOW WOULD AN IDEAL CODE LOOK?

```
class MyDuck {  
public:  
    int x;  
    MyDuck(int x) : x(x) {}  
    std::string do_quack() const { return "quack"; }  
};  
  
auto a = make_virtual_any<int>(7);  
std::cout << "Printing " << a << std::endl;  
  
a = make_virtual_any<MyDuck>(MyDuck(3));  
std::cout << "My obj " << a << " has .x == " << a.attr("x")  
    << ", .do_quack() == " << a.attr("do_quack")() << std::endl;
```


WE CAN DO THIS!



HOLDING CLASS : VIRTUAL_ANY

```
class virtual_any_interface;
class virtual_any {
    std::shared_ptr<virtual_any_interface> _impl;
public:
    virtual_any(std::shared_ptr<virtual_any_interface> elem) : _impl(elem) {}
    virtual_any attr(const std::string& name);
    virtual_any operator()();
    friend std::ostream& operator<<(std::ostream& os, const virtual_any& self);
};

class virtual_any_interface {
public:
    virtual virtual_any attr(const std::string& name) = 0;
    virtual virtual_any call() = 0;
    virtual std::ostream& stream(std::ostream& os) const = 0;
};
```

VIRTUAL_ANY_INTERFACE

- `virtual_any` calls a virtual method on the held type
 - `virtual_any_interface`
- We implement `virtual_any_interface` for each type
 - Using reflection!
- Handle some special function operators natively
 - `operator<<`
 - `operator()`
 - `operator+`
 - Can handle more, but need to figure out how to handle variadic args

IMPLEMENTATION CLASS : VIRTUAL_ANY_IMPL

```
template <typename T>
class virtual_any_impl : public virtual_any_interface {
    T _value;
    constexpr std::vector<std::pair<std::string, virtual_any>> get_attrs(); // All members and methods of T
public:
    virtual_any_impl(T& value) : _value(std::forward<T>(value)) {}
    virtual virtual_any attr(const std::string& name) override;
    virtual virtual_any call() override; // .attr could've returned a std::function<virtual_any(void)>
    virtual std::ostream& stream(std::ostream& os) const override;
};
```

VIRTUAL_ANY_IMPL: IGNORE SPECIAL CASES

```
template <typename T>
constexpr std::vector<std::pair<std::string, virtual_any>> virtual_any_impl<T>::get_attrs() {
    using T2 = std::remove_cvref_t<T>;
    if constexpr(!meta::test_type (^std::is_class_v, ^T)) {
        return {};
    } else if constexpr(std::is_same_v<T2, std::string> ||
        std::is_same_v<T2, std::function<virtual_any(void)>>) {
        return {};
    } else {
        std::vector<std::pair<std::string, virtual_any>> attrs;
        // Actual logic
        return attrs;
    }
}
```

VIRTUAL_ANY_IMPL: GET ALL MEMBERS AND METHODS

```
template <typename T>
constexpr std::vector<std::pair<std::string, virtual_any>> virtual_any_impl<T>::get_attrs() {
    [:expand(meta::members_of(^T)):] >> [&]<auto e> {
        if constexpr(!meta::is_public(e)) return;
        else {
            auto name = std::string(std::meta::name_of(e));
            if constexpr(meta::is_nonstatic_data_member(e)) {
                attrs.push_back({name, make_virtual_any(_value.[:e:])});
            } else if constexpr(meta::is_function(e) && !meta::is_constructor(e) && !meta::is_destructor(e) &&
                                !meta::is_special_member(e)) {
                std::function<virtual_any(void)> l = [this]() { return make_virtual_any(_value.[:e:]()); };
                attrs.push_back({name, make_virtual_any(l)});
            }
        }
    };
}
```

VIRTUAL_ANY_IMPL: IMPLEMENT THE `ATTR` METHOD

```
template <typename T>
virtual_any virtual_any_impl<T>::attr(const std::string& name) {
    auto attrs = get_attrs();
    if (attrs.size() == 0) {
        throw std::runtime_error("No attributes found");
    }
    for (const auto& [name2, e] : attrs) {
        if (name2 == name) {
            return e;
        }
    }
    std::stringstream ss;
    ss << "Attribute " << name << " not found in object of type " << std::meta::name_of(^T);
    throw std::runtime_error(ss.str());
}
```

VIRTUAL_ANY_IMPL: IMPLEMENT OTHER HELPER METHODS

```
template <typename T> virtual_any virtual_any_impl<T>::call() {
    if constexpr(std::is_same_v<T, std::function<virtual_any(void)>>) {
        return _value();
    } else {
        std::stringstream ss;
        ss << "Object of type " << std::meta::name_of(^T) << " is not callable";
        throw std::runtime_error(ss.str());
    }
}

template <typename T> std::ostream& virtual_any_impl<T>::stream(std::ostream& os) const {
    if constexpr(is_streamable<std::stringstream, T>::value) {
        os << _value;
    } else {
        os << "(non-streamable object of type " << std::meta::name_of(^T) << ")";
    }
    return os;
}
```

AND THERE IT IS!

IT WORKS :)

```
int main() {  
    auto a = make_virtual_any<int>(7);  
    std::cout << "Printing " << a << std::endl;  
  
    a = make_virtual_any<MyDuck>(MyDuck(3));  
    std::cout << "My obj " << a << " has .x == " << a.attr("x") << std::endl;  
    std::cout << "My obj " << a << " has .do_quack() == " << a.attr("do_quack")() << std::endl;  
}
```

```
$ ./virtual_any/main.out
```

```
Printing 7
```

```
My obj MyDuck(3) has .x == 3
```

```
My obj MyDuck(3) has .do_quack() == quack quack quack
```

DISCUSSION


- Name is solid I know ;)
- Slow to run
 - Not the worst thing for writing script like code from time to time
 - Especially useful when dealing with containers and dictionaries and lists like in pythonic code
- Slow to compile
 - This one's actually sad and hard to get around
 - Compilation can't even be shared across compilation units
 - Unless... Language idioms

WHAT DO WE GET OUT OF THIS?

- A fun new experiment with Reflection?
- Offers a clean duck-typed syntax, setup at compile time
 - Could start implementing better DSLs in C++?
- Removes need for complex ABI compatibility at module boundary
 - Sound familiar?

ALTERNATIVES TO REFLECTION

How did we survive
without this till now?



HUMAN REFLECTION

- Well, we've been writing boilerplate
 - Python bindings
 - Serialization / deserialization / CLI parsers
- Our code, and our libraries are worse off
- Compile-time reflection only makes for cleaner APIs
 - There's no “new” semantic thing we can do now that we couldn't do
 - We *can* do “new” syntactic things to avoid redundancy in code

SOURCEGEN!

- Protobuf
 - (in)famous library loved and hated by everyone at the same time
 - Lets you do some primitive “reflection-like” stuff
 - Enum \Leftrightarrow String
- Apache Avro
 - Sourcegen with a “handshake” / “header”
 - Better “reflection-like” API than protobuf

SOURCEGEN! FT. PROTOBUF

```
enum VehicleType {  
    UNDEFINED = 0;  
    CAR = 1;  
    TRUCK = 2;  
    MOTORCYCLE = 3;  
}  
  
message Vehicle {  
    string model = 1;  
    VehicleType type = 2;  
    repeated string features = 3;  
}
```

```
enum VehicleType { UNDEFINED = 0, CAR, TRUCK, MOTORCYCLE };  
bool VehicleType_IsValid(int value);  
const ::google::protobuf::EnumDescriptor* VehicleType_descriptor();  
inline const ::std::string& VehicleType_Name(VehicleType value);  
inline bool VehicleType_Parse(const ::std::string& name, VehicleType* value);  
  
class Vehicle {  
public:  
    const std::string& model() const;  
    void set_model(const std::string& value); ...  
    const ::google::protobuf::RepeatedPtrField<std::string>& features() const;  
private:  
    std::string model_;  
    VehicleType type_;  
    ::google::protobuf::RepeatedPtrField<std::string> features_;  
};
```

SOURCEGEN! FT. AVRO

```
{
  "namespace": "example",
  "type": "record",
  "name": "Vehicle",
  "fields": [
    {"name": "model", "type": "string"},
    {"name": "type",
     "type": {
       "type": "enum",
       "name": "VehicleType",
       "symbols": [...]
     }
    },
    {"name": "features",
     "type": {
       "type": "array", "items": "string"
     }
    }
  ]
}
```

```
struct Vehicle {
    std::string model;
    VehicleType type;
    std::vector<std::string> features;
    void serialize(std::ostream& out) const {
        auto schema = avro::compileJsonSchemaFromString(schemaJson);
        avro::EncoderPtr encoder = avro::binaryEncoder();
        encoder->init(out);
        avro::encode(*encoder, *this);
    }
    void deserialize(std::istream& in) {
        auto schema = avro::compileJsonSchemaFromString(schemaJson);
        avro::DecoderPtr decoder = avro::binaryDecoder();
        decoder->init(in);
        avro::decode(*decoder, *this);
    }
};
```


SOURCEGEN! FT. LLM

```
struct ConfigOptions {  
    bool enable;  
    int count;  
    std::string mode;  
};  
  
int main() {  
    boost::program_options po;  
    po.add_options()  
        ("enable", boost::program_options::value<bool>(), "enable")  
        ("count", boost::program_options::value<int>(), "count")  
        ("mode", boost::program_options::value<std::string>(), "mode");  
}
```

AMONG OTHERS...

- Classdesc
 - Another form of sourcegen
- <https://github.com/boost-ext/reflect>
 - Template metaprogramming to its limits

[:FIN:]

Saksham Sharma

